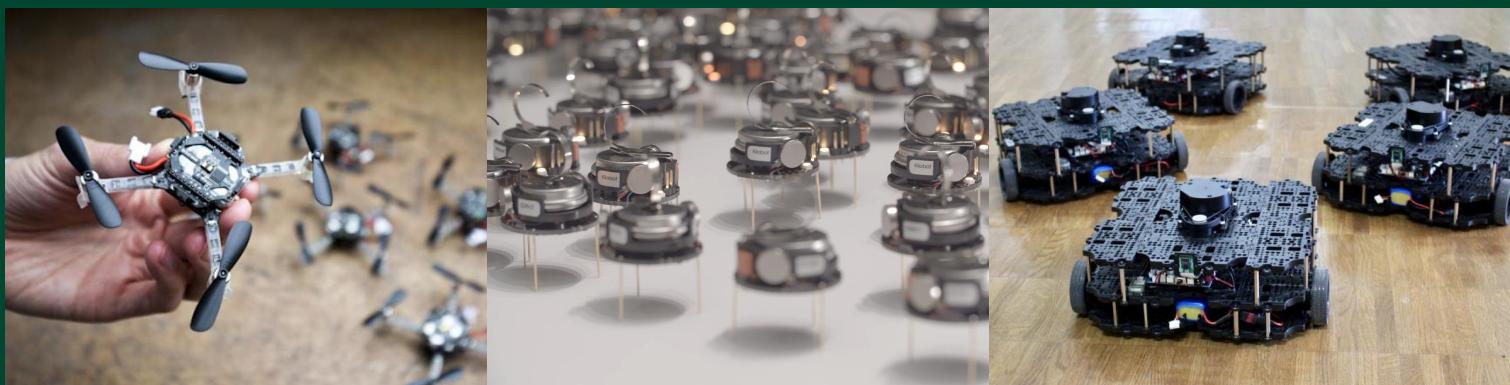


ECE693H, Spring 2025: Multi-robot System Design

“Multi-robot Systems 2”



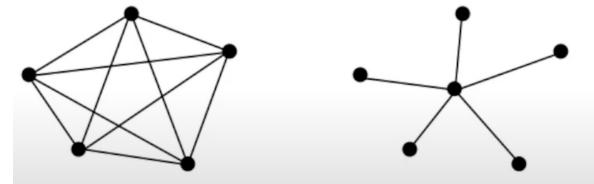
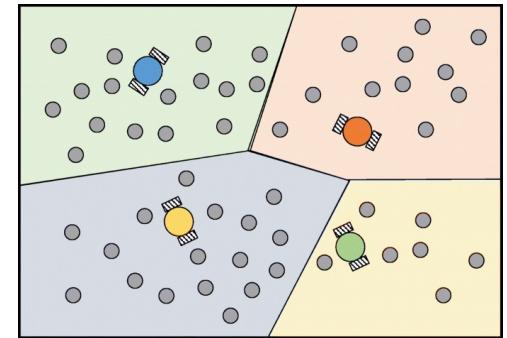
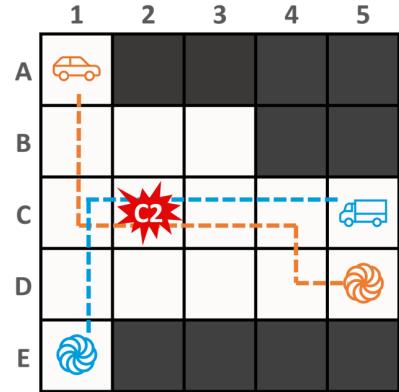
Dr. Daniel Drew



MRS Survey Recap



A group (>1) of robots with a common objective



Questions?

Class Overview

1. Lecture (~20 minutes)
2. Astound the Class (~40 minutes)

Multi-robot Systems: Lecture 2 of 4

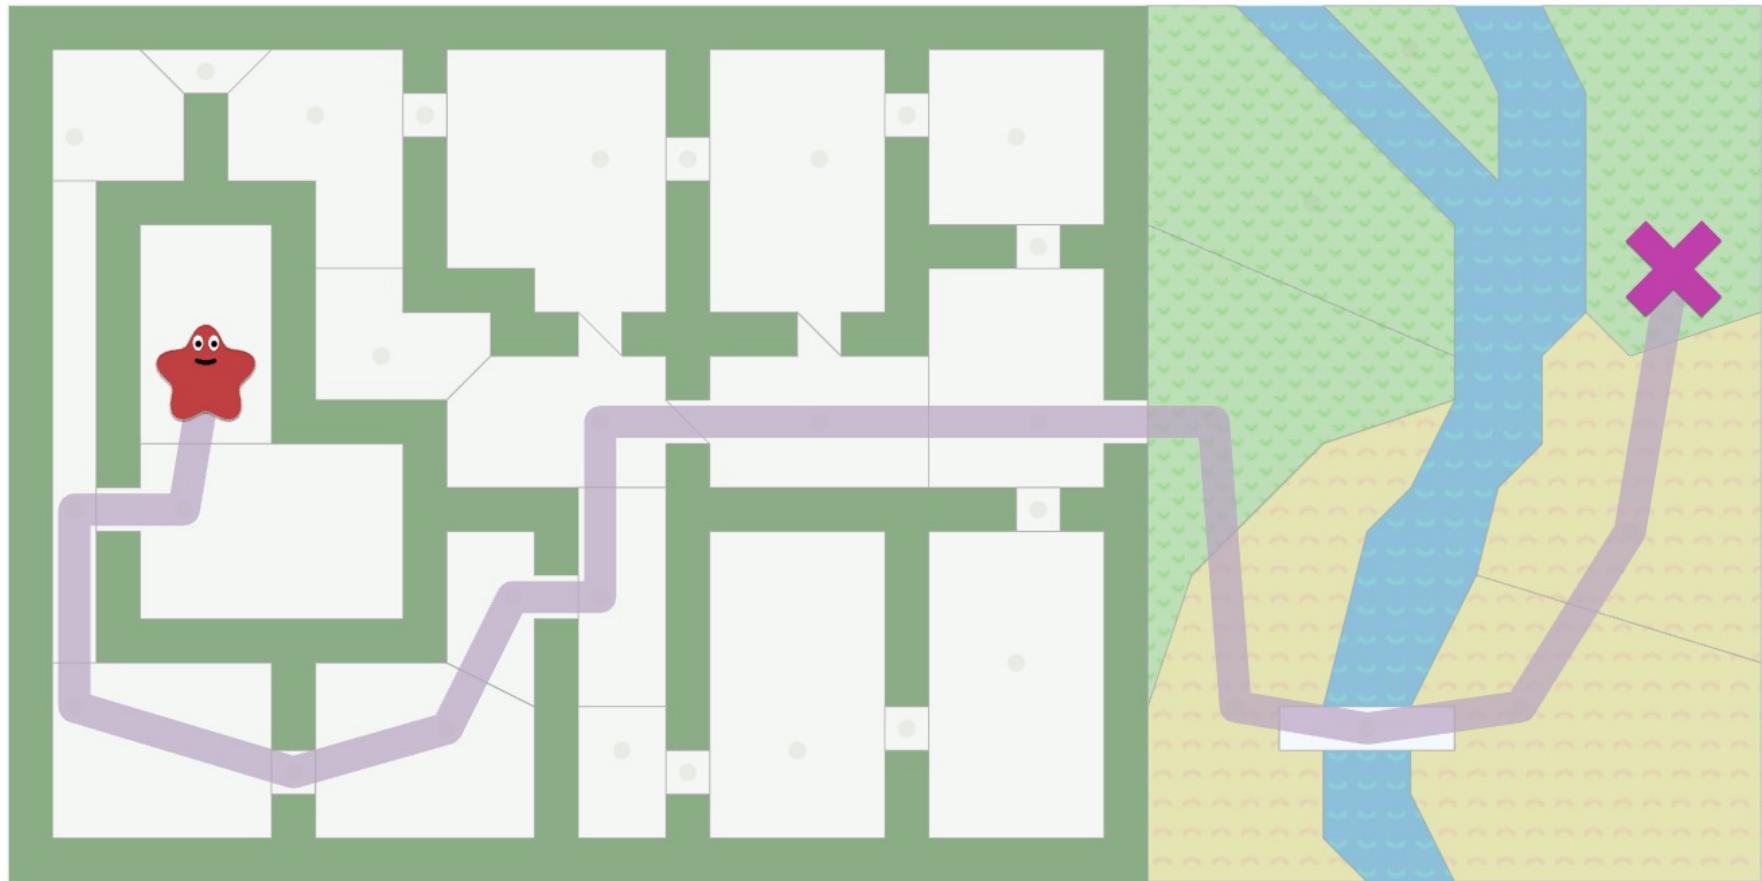
Lecture 1: MRS Survey

Lecture 2: Motion Planning

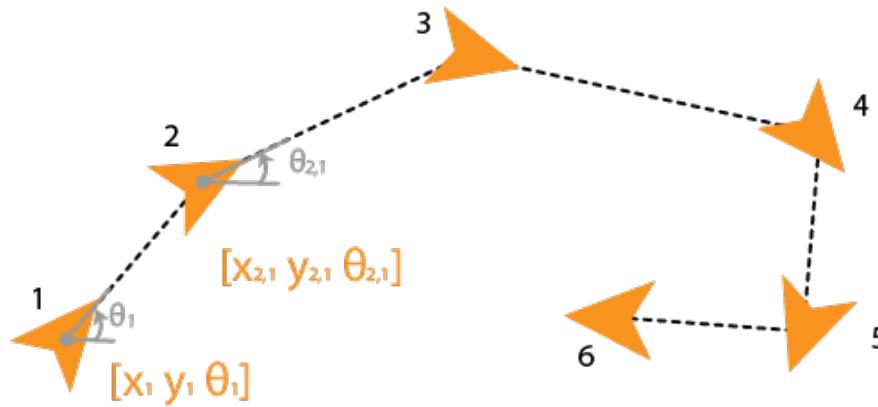
Lecture 3: Task Assignment (+Motion Planning if needed)

Lecture 3: Communication and Control Paradigms

The Motion Planning Problem



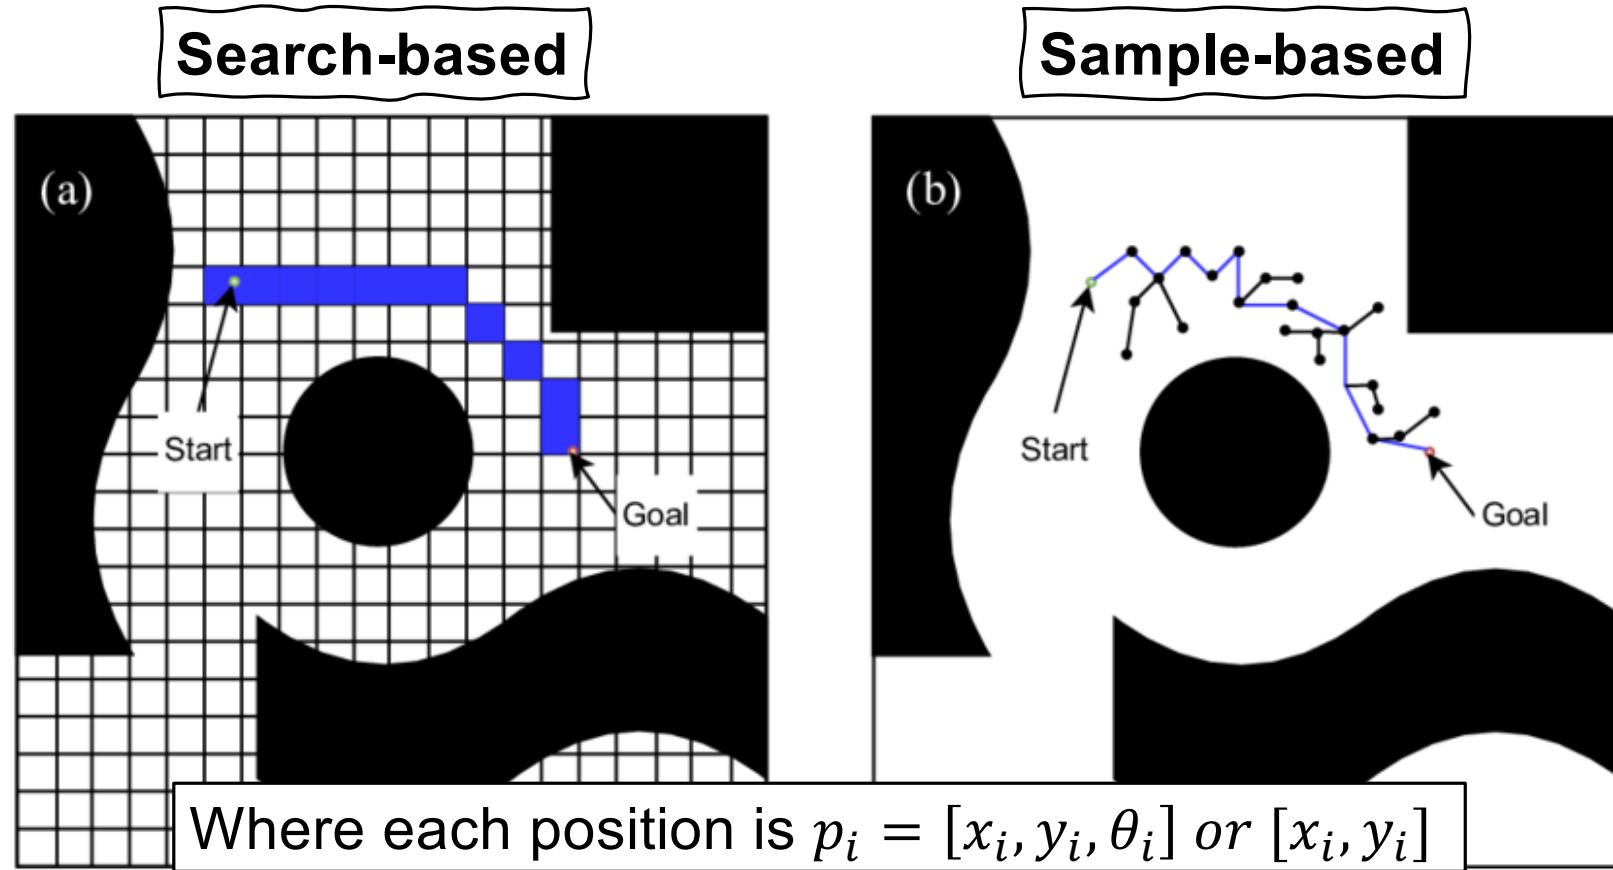
The Motion Planning Problem



What set of poses best gets the robot from A to B?



Review of Single-Robot Motion Planning



Multi-robot Motion Planning: Taxonomies

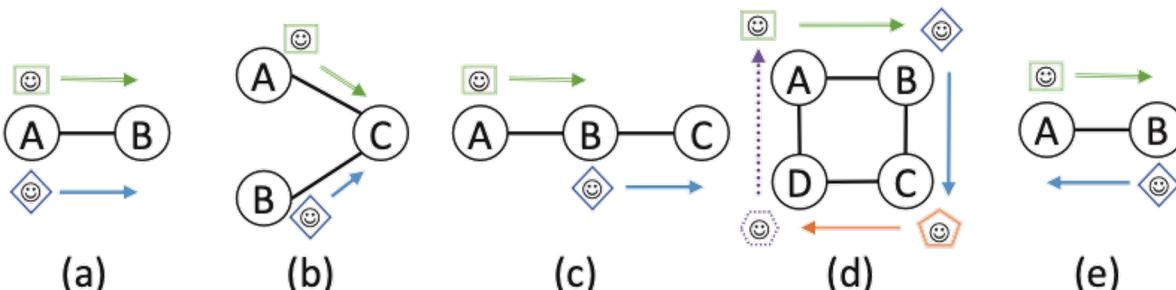
Search-based versus Sample-based

Same dichotomy!

Coupled versus Decoupled

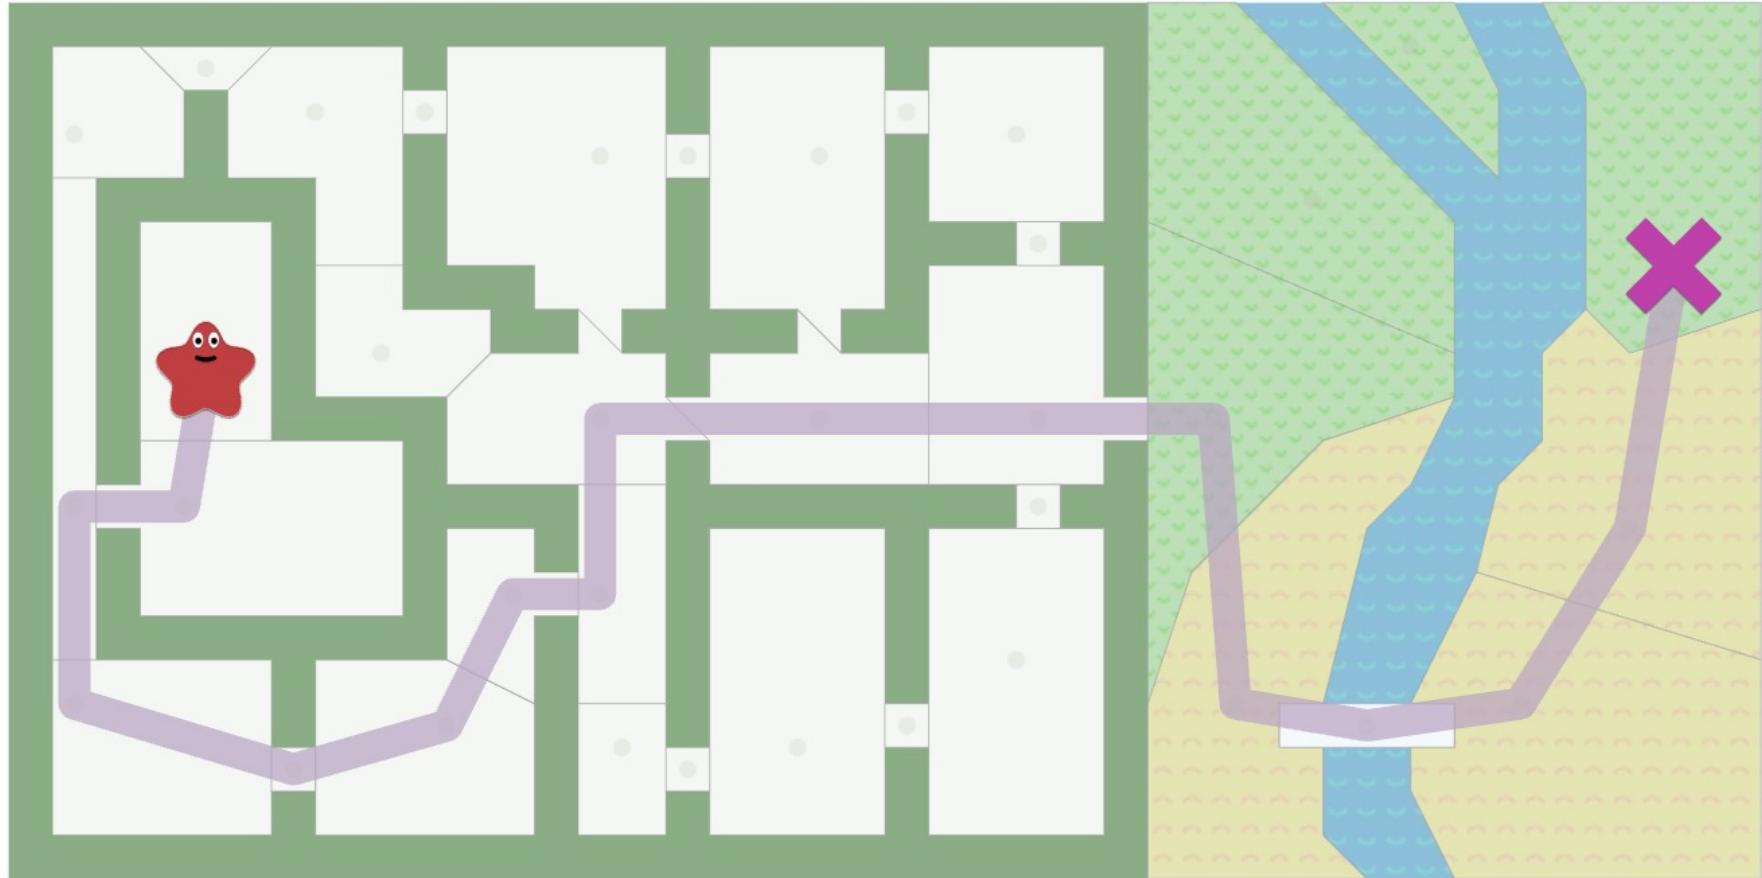
Plan in the full (“global”) joint configuration space versus plan in individual (“local”) configuration spaces

Conflicts

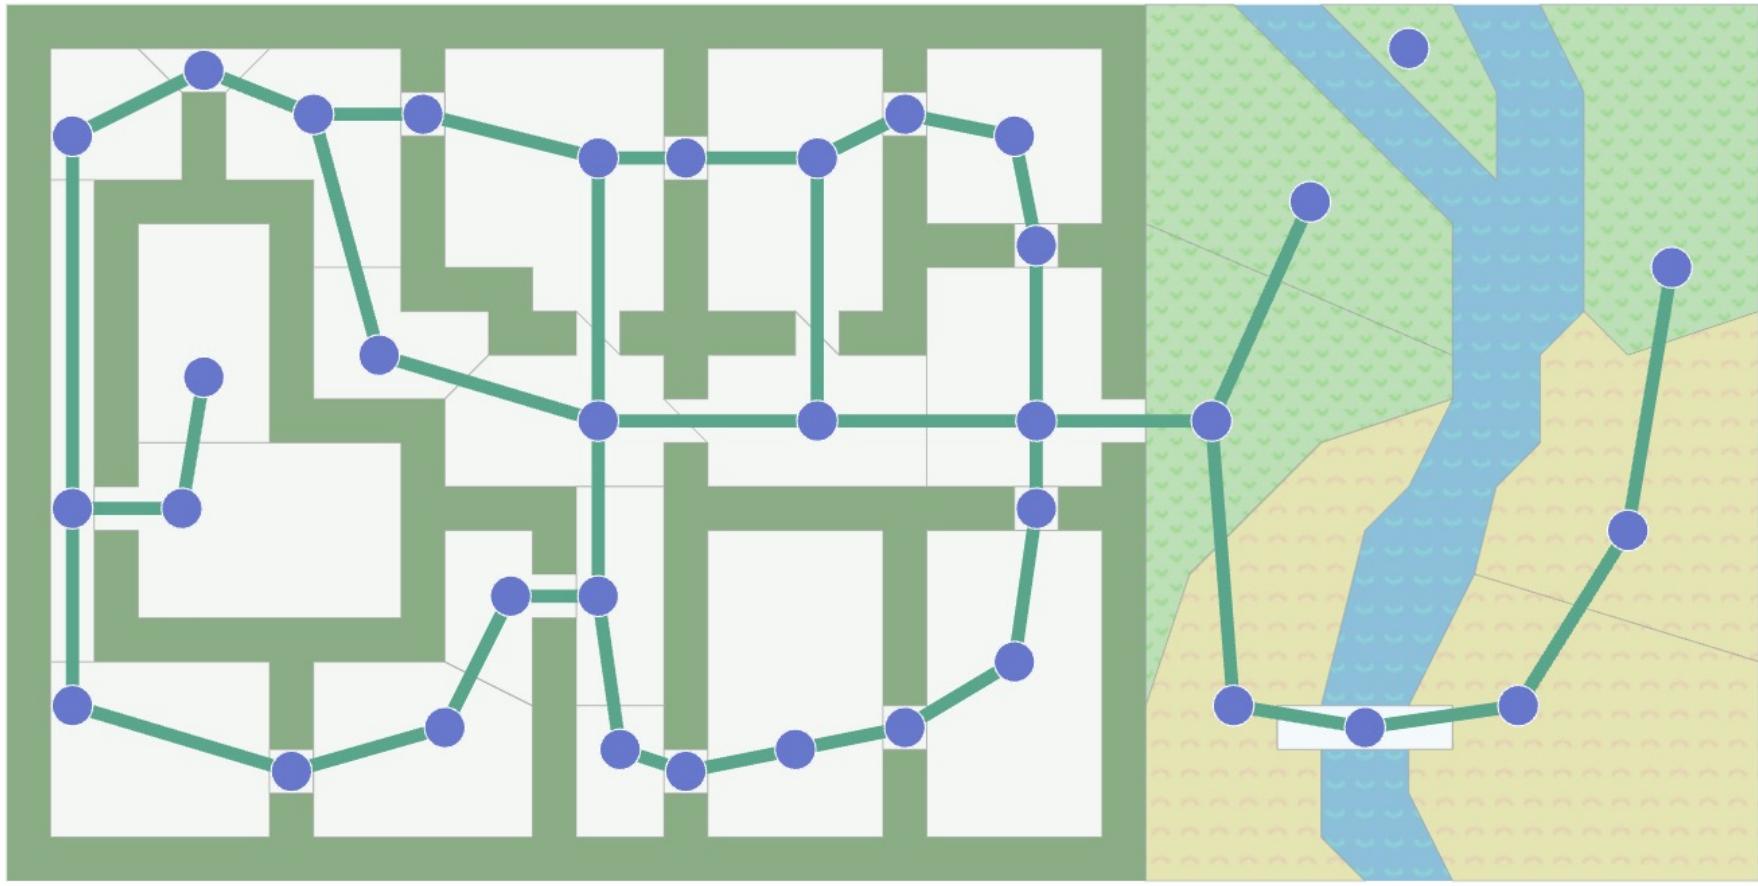


(a) is an edge conflict, (b) a vertex conflict, (c) a following conflict, (d) a cycle conflict, and (e) a swapping conflict.

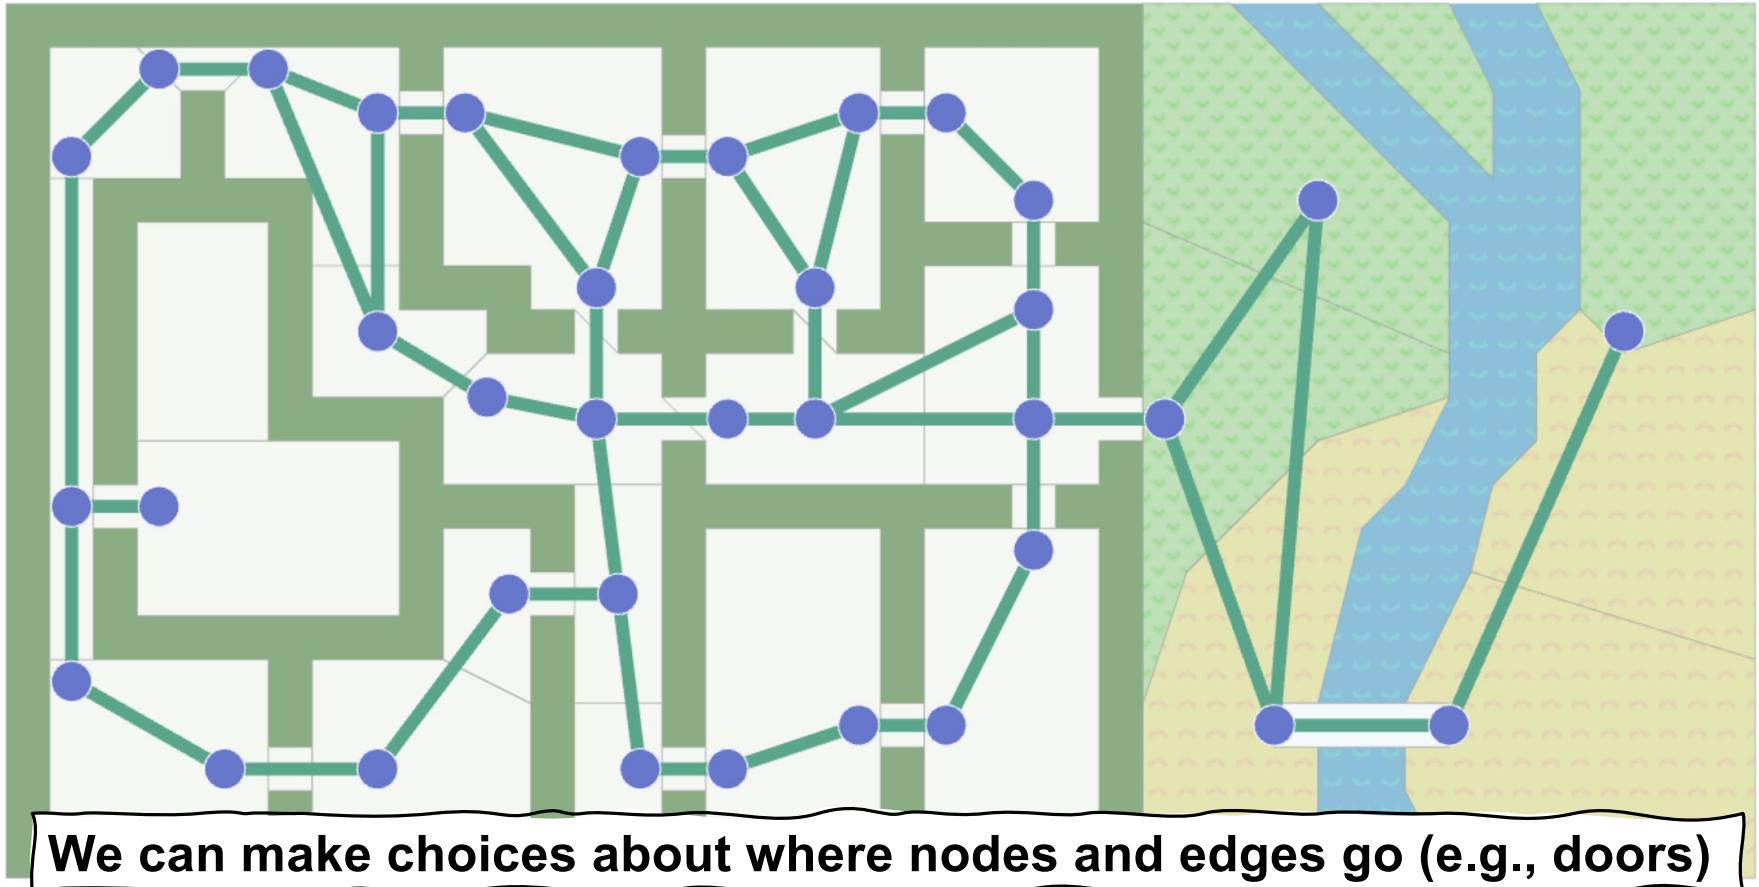
Single Robot Planning: A*



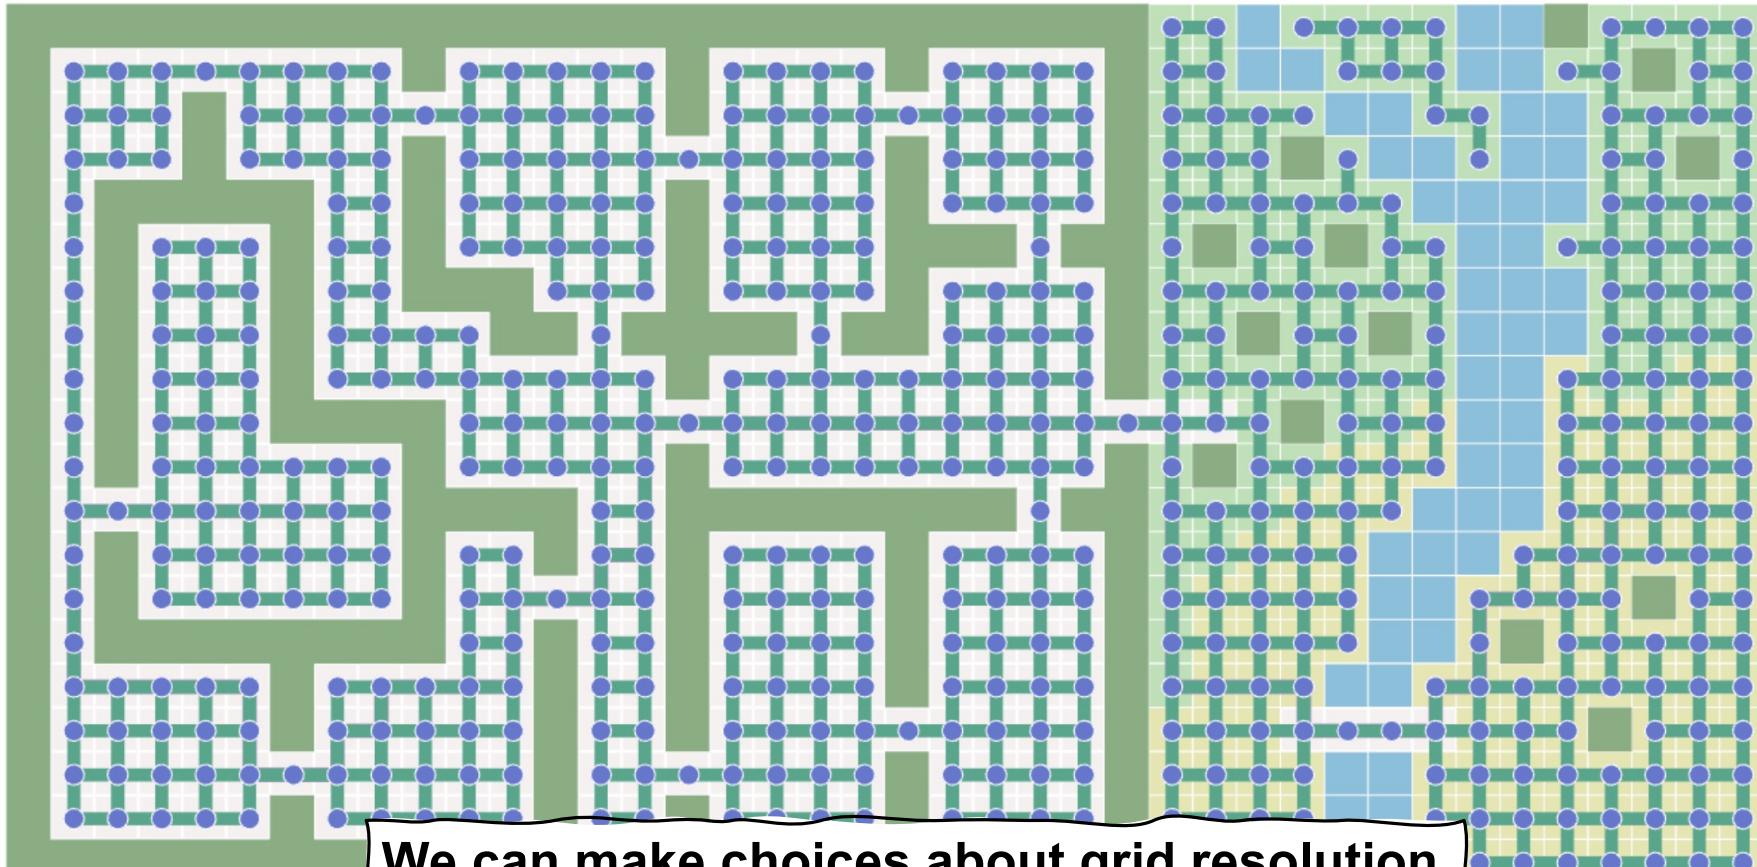
Single Robot Planning: A*



Single Robot Planning: A*



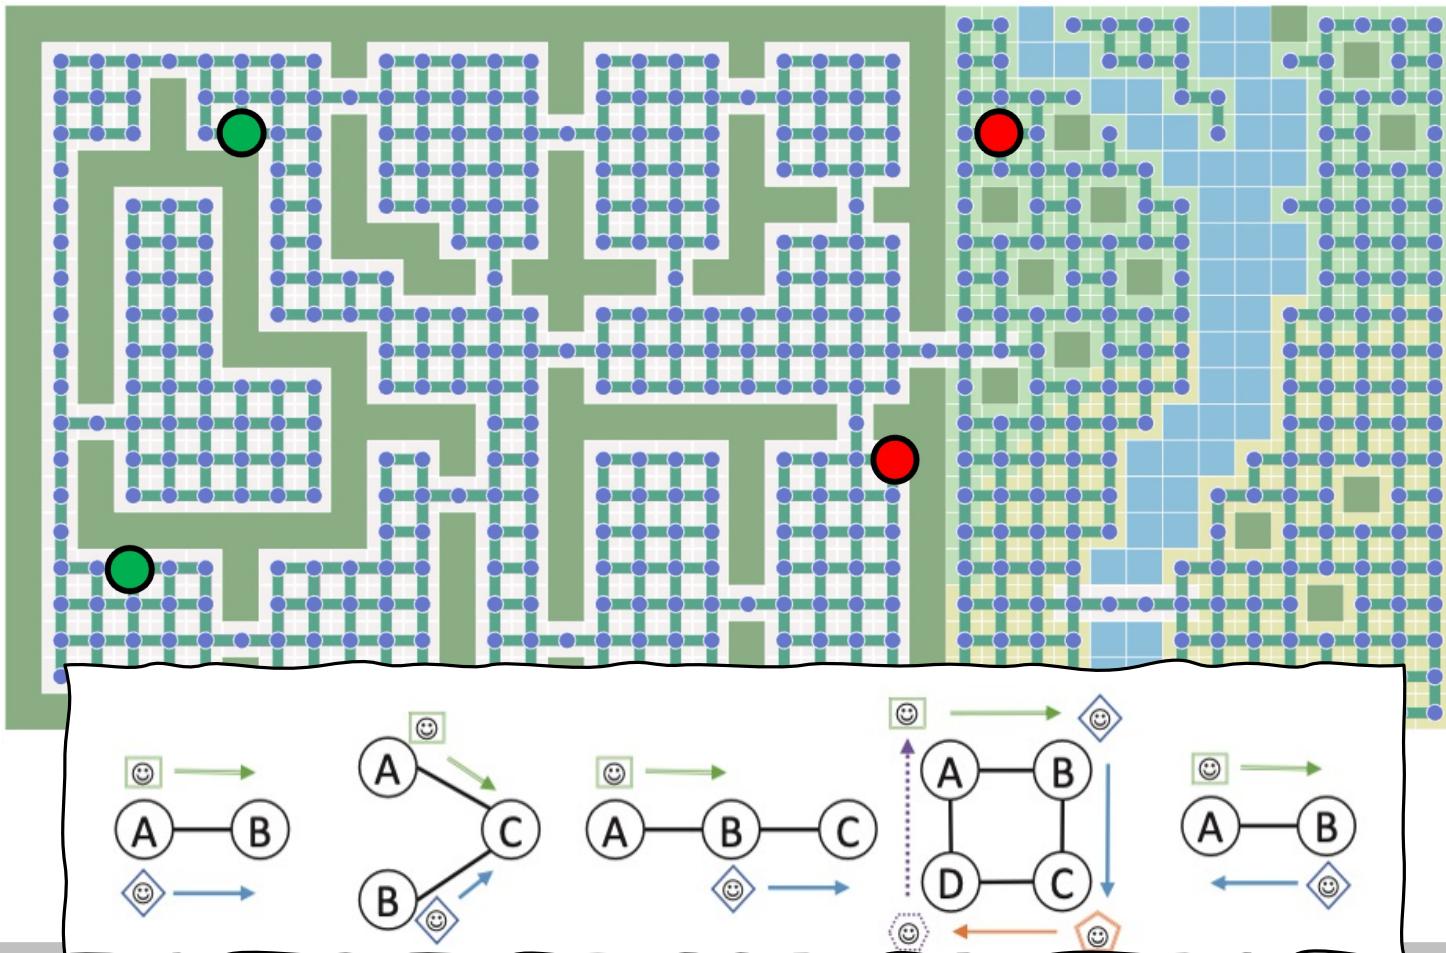
Single Robot Planning: A*



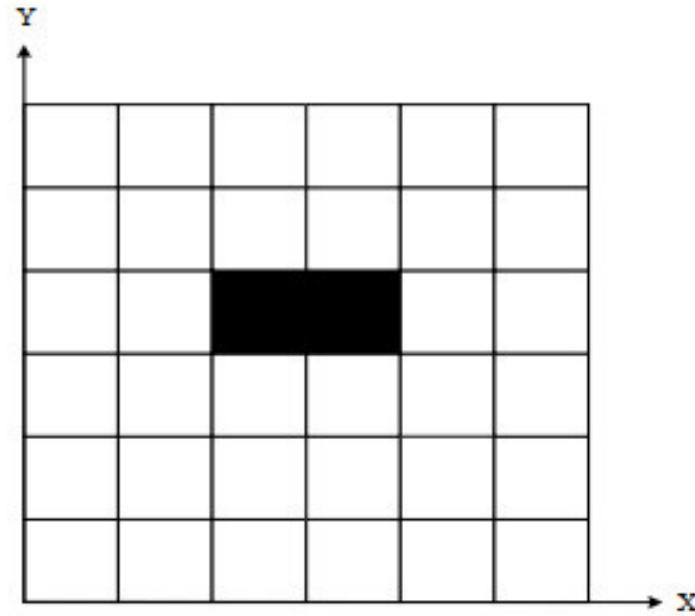
Single Robot Planning: A*

[https://www.redblobgames.com/pathfinding/
a-star/introduction.html](https://www.redblobgames.com/pathfinding/a-star/introduction.html)

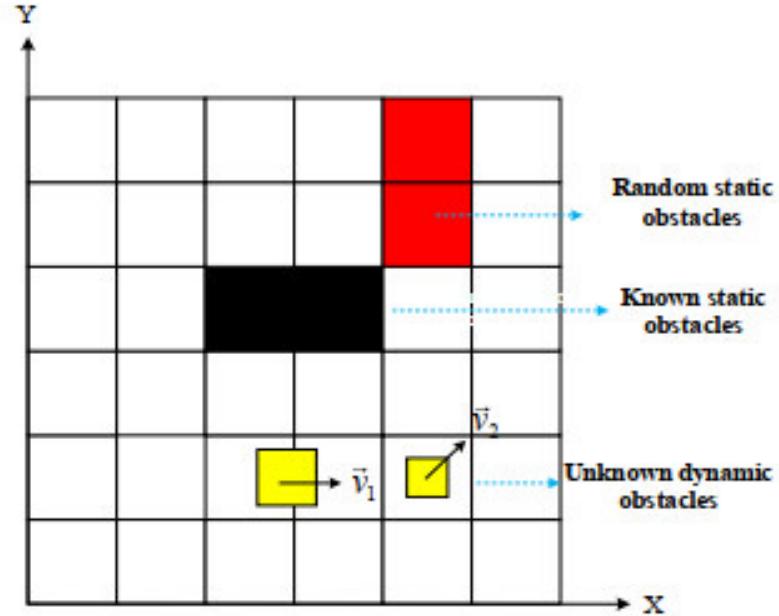
Search-based Methods for MRS



A* with Replanning: Local Repair (LRA*)

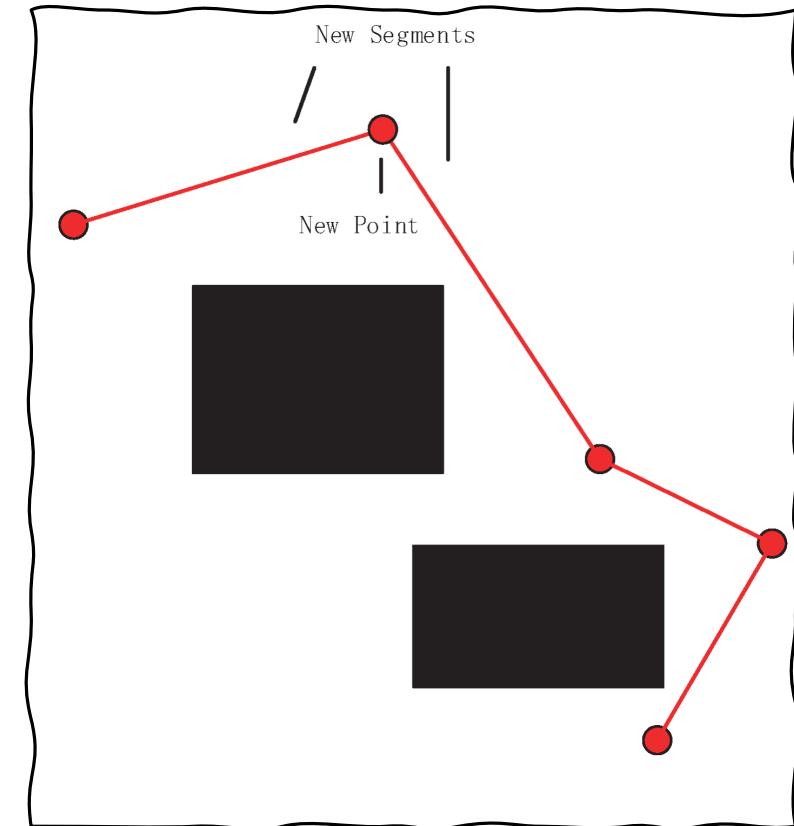
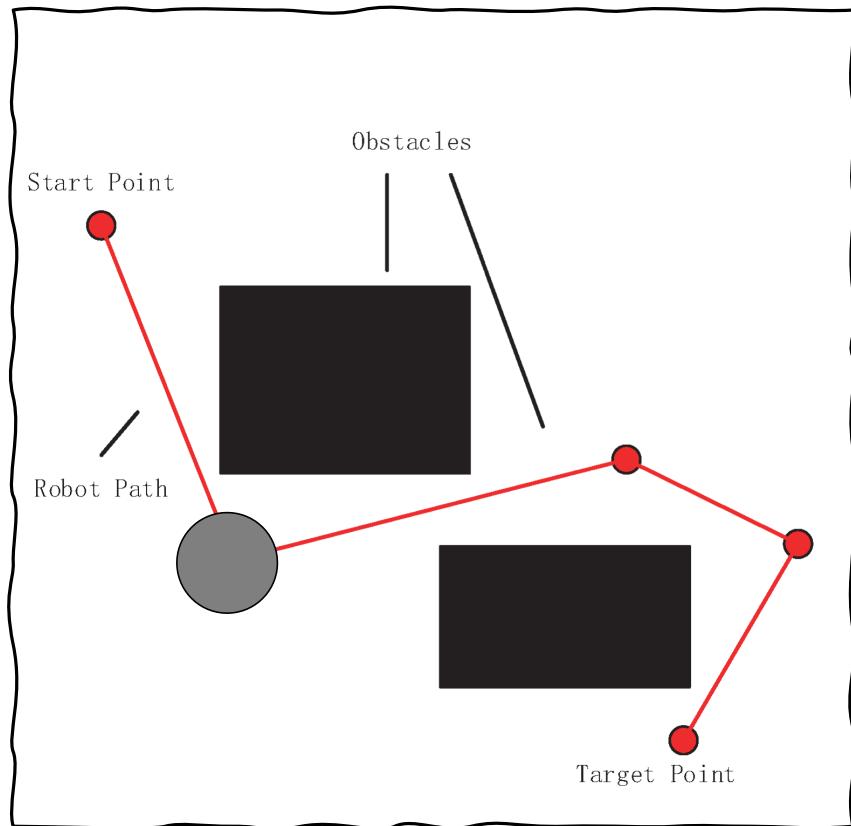


(a) Traditional global static environment



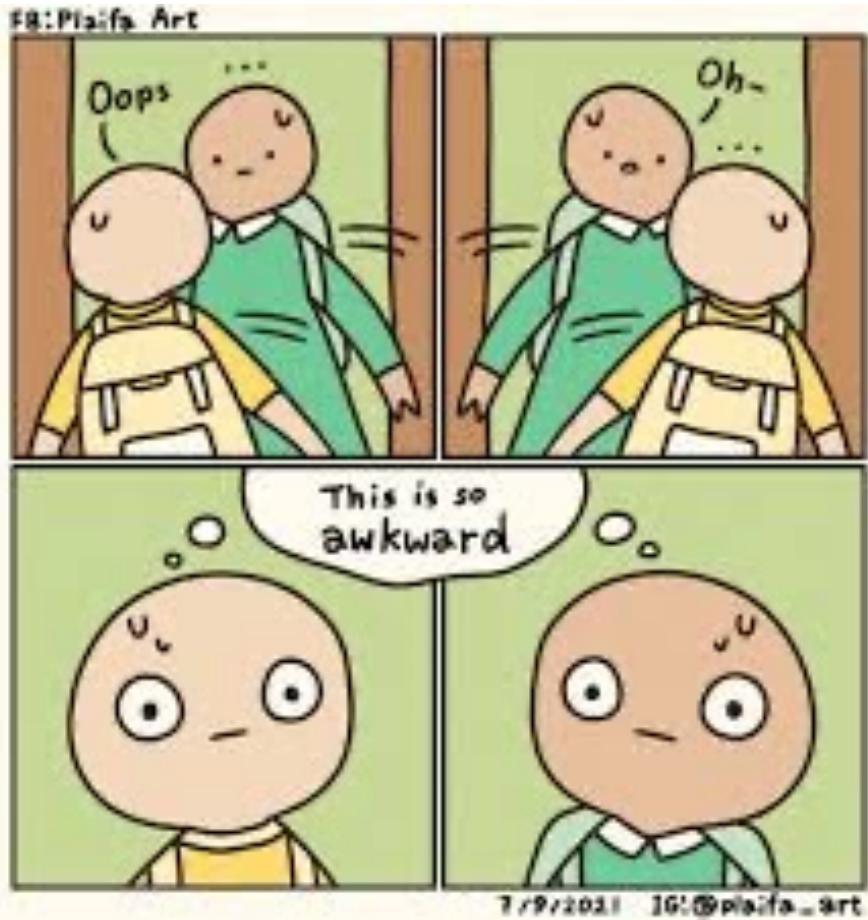
(b) Global static + local dynamic environment in this paper

A* with Replanning: Local Repair (LRA*)



Don't need to replan the entire route

A* with Replanning: Local Repair (LRA*)



Local replanning without additional rules leads to **collision cycles**.

One workaround: “**agitation**”: after a cycle, add random noise to the graph edges.

Cooperative A* (CA*)

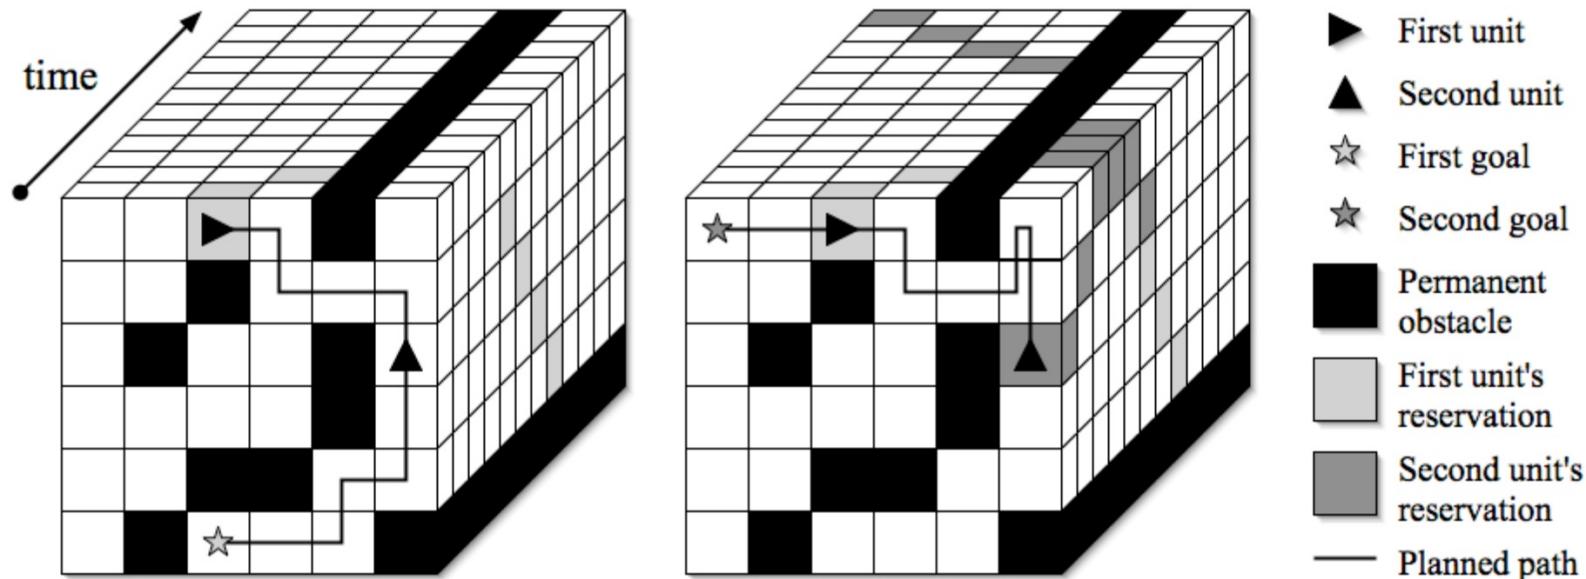


Figure 1 Two units pathfinding cooperatively. (A) The first unit searches for a path and marks it into the reservation table. (B) The second unit searches for a path, taking account of existing reservations, and also marks it into the reservation table.

$[x_i, y_i] \rightarrow [x_i, y_i, t_i]$

Add command *wait*

Challenge: curse of dimensionality

Conflict Based Search (CBS)

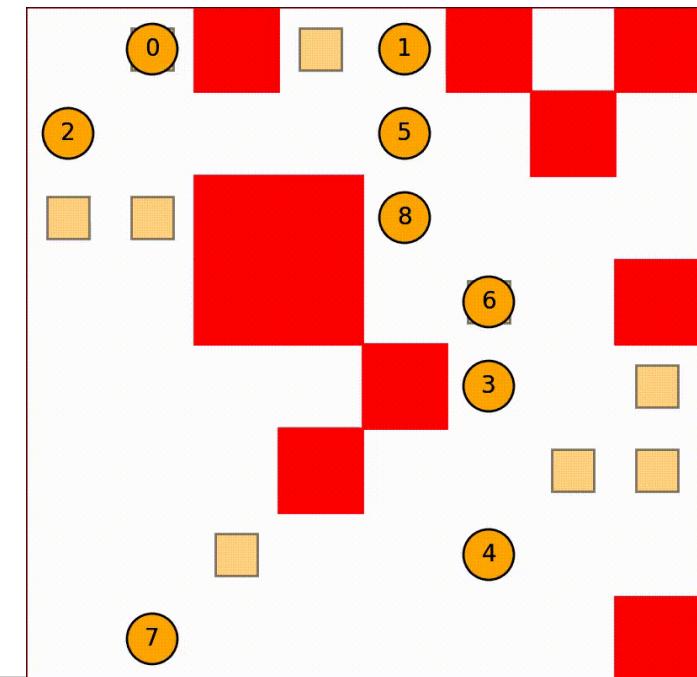
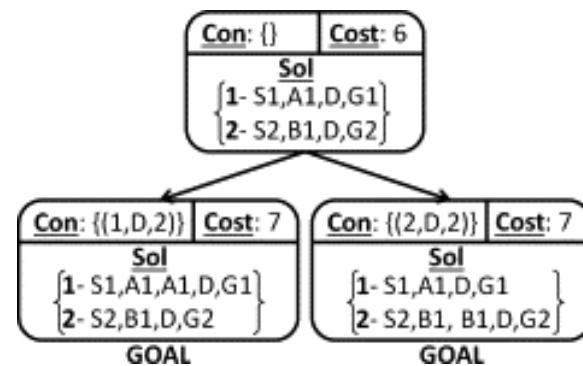
Input: MAPF instance

```

1 Root.constraints =  $\emptyset$ 
2 Root.solution = find individual paths by the low level()
3 Root.cost = SIC(Root.solution)
4 insert Root to OPEN
5 while OPEN not empty do
6   P  $\leftarrow$  best node from OPEN // lowest solution cost
7   Validate the paths in P until a conflict occurs.
8   if P has no conflict then
9     return P.solution // P is goal
10  C  $\leftarrow$  first conflict  $(a_i, a_j, v, t)$  in P
11  [REDACTED]
12  [REDACTED]
13  [REDACTED]
14  [REDACTED]
15  [REDACTED]
16  [REDACTED]
17  [REDACTED]
18  [REDACTED]
19  foreach agent  $a_i$  in C do
20    A  $\leftarrow$  new node
21    A.constraints  $\leftarrow$  P.constraints +  $(a_i, v, t)$ 
22    A.solution  $\leftarrow$  P.solution
23    Update A.solution by invoking low level( $a_i$ )
24    A.cost = SIC(A.solution)
25    if A.cost <  $\infty$  // A solution was found then
      | Insert A to OPEN
  
```

Conflict (a_i, a_j, v, t) (agent 1 and 2 occupy vertex at time)

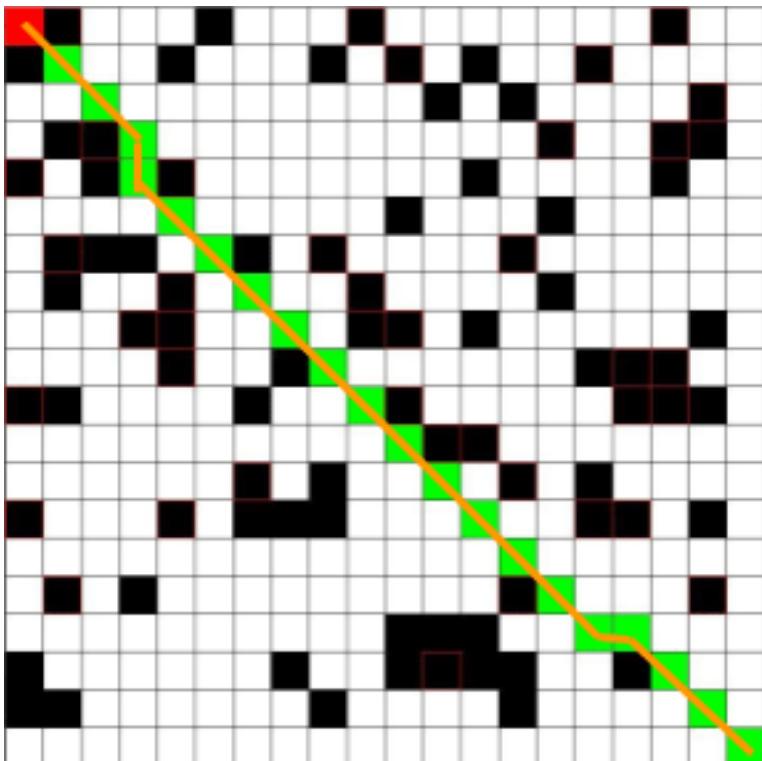
Constraint (a_i, v, t) (agent can't occupy vertex at time)



A* PATHFINDING

EPISODE 01

D* and D* Lite



Two separate lists maintained:

“**G**” score = Cost to connect the starting point to the current point

“**RHS**” score = Connection cost from current node to next node, inf. for obstacle; aka one-step lookahead cost

→when $G(s) \leq RHS(s)$, queue for update.

Terminology:

s = current node

s' = predecessor node

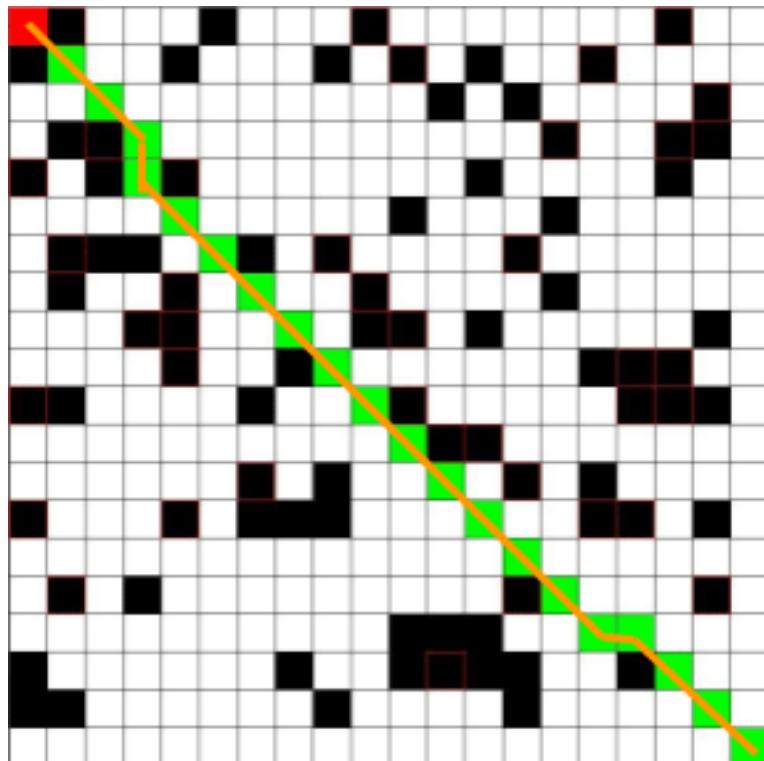
$C(s', s)$ = edge cost for $s' \rightarrow s$

$G(s)$ = current cost to arrive at that node

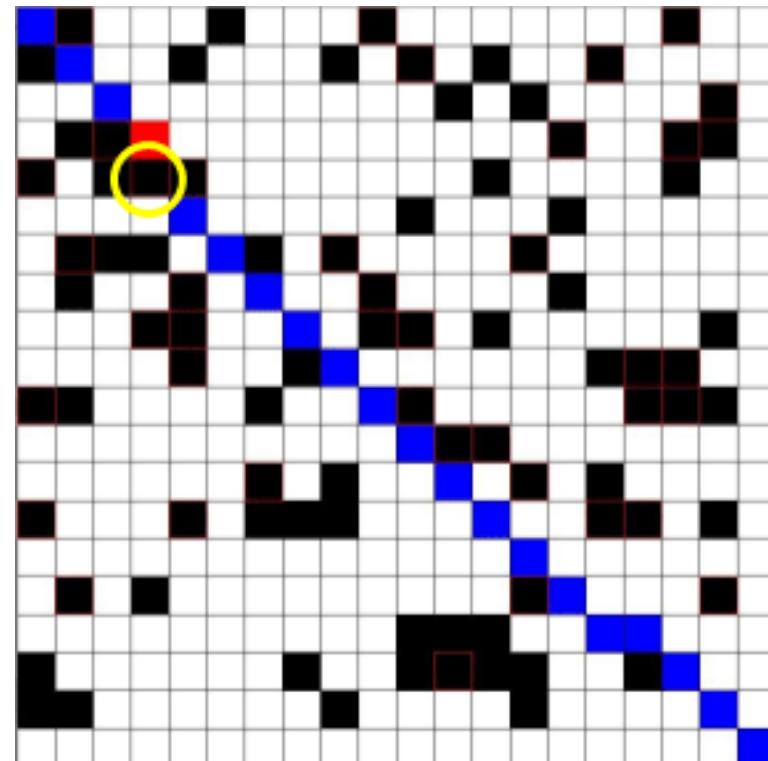
$G(s) = G(s') + C(s', s)$

$RHS(s) = \min \{G(s') + C(s', s)\}$

D* and D* Lite

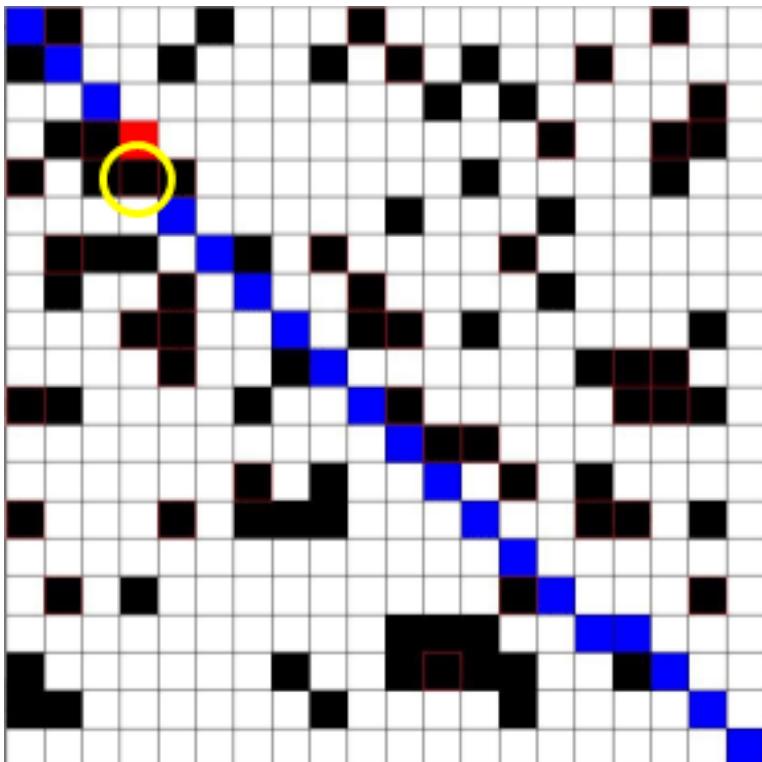


Initial plan

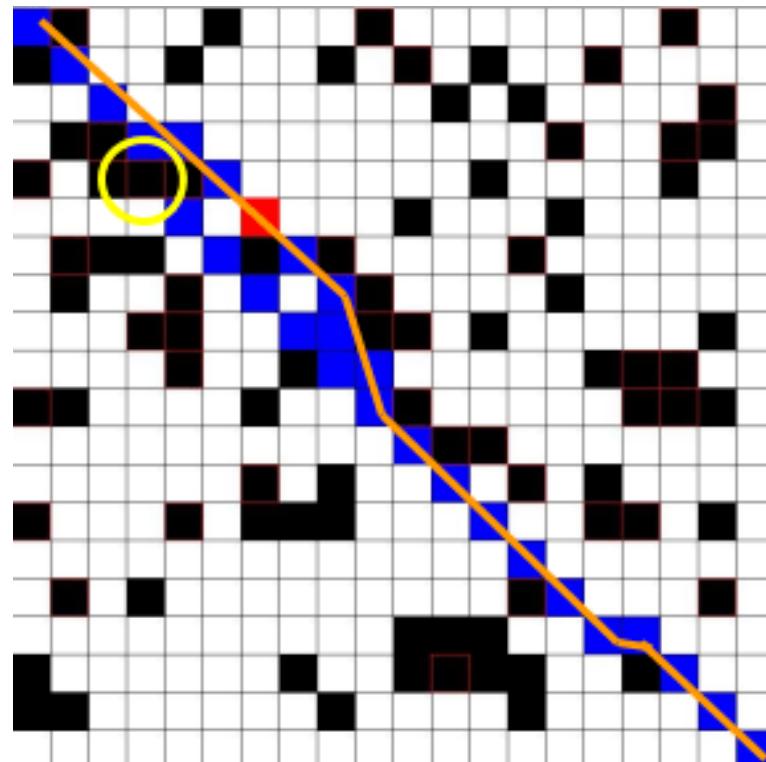


Obstacle encountered

D* and D* Lite

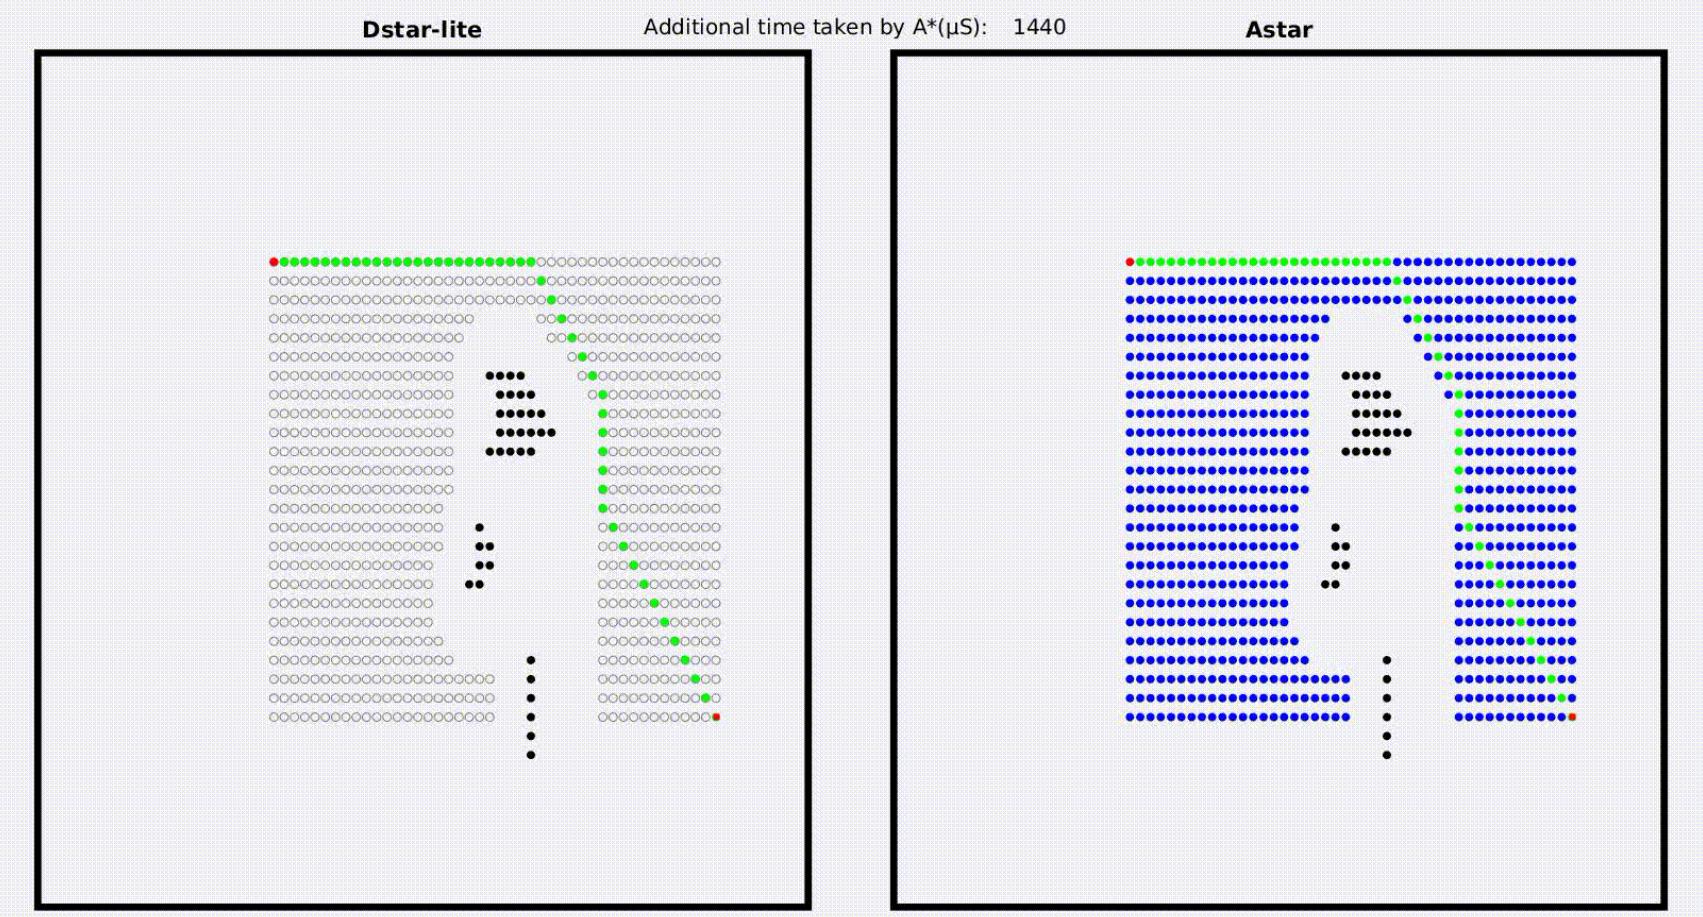


Obstacle encountered



Plan updated

D* and D* Lite



Contrasting These Approaches

A* with replanning:

Initial planning: A*

Replanning: Re-run A*

Efficiency: Poor

Optimality: High

D* Lite:

Initial planning: A*-ish in reverse

Replanning: Update lookahead cost RHS(s)

Efficiency: High

Optimality: Variable

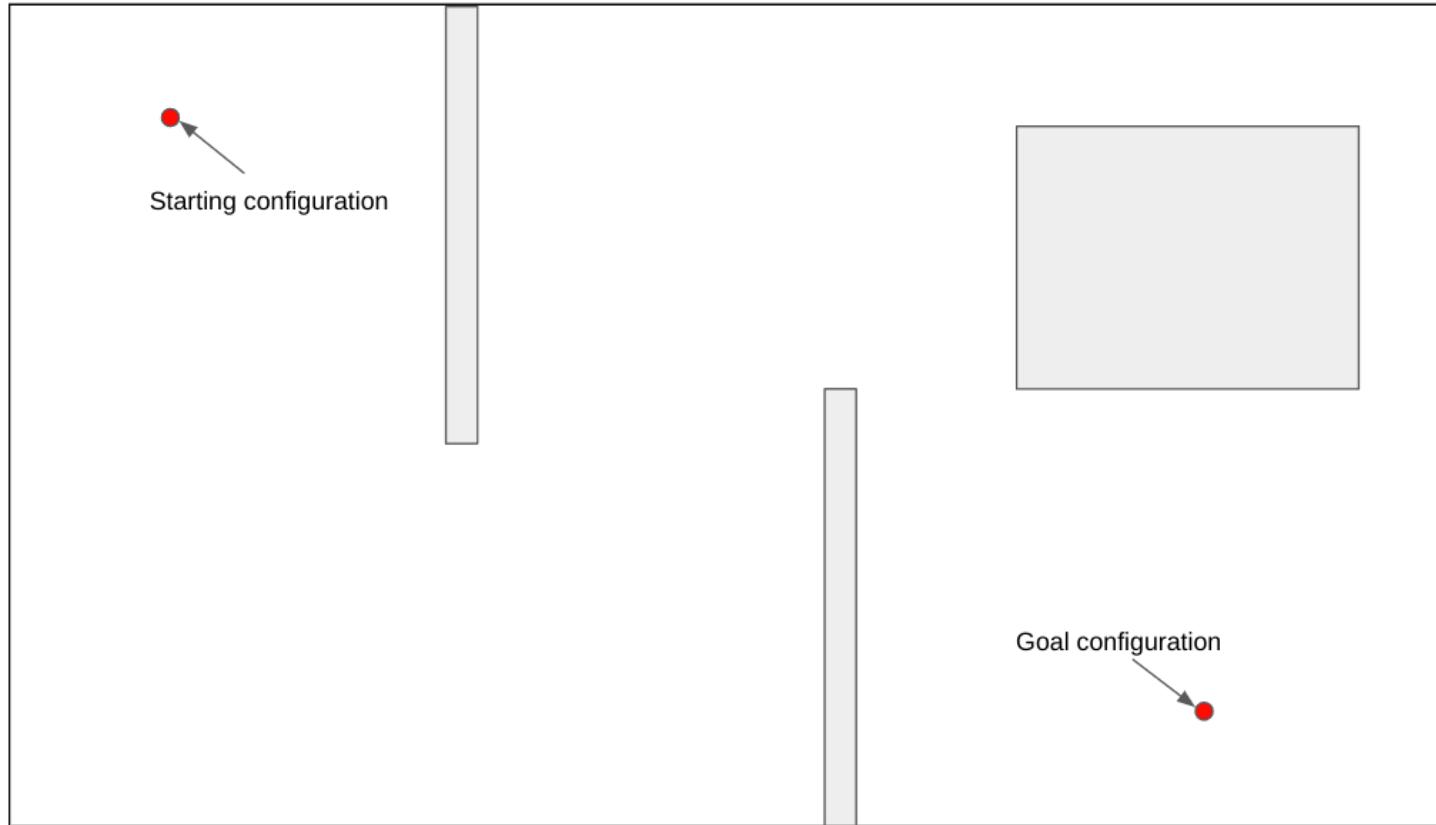
D* Lite for mostly dynamic environments, A* for mostly static

Single Robot Planning: Sampling-based Methods

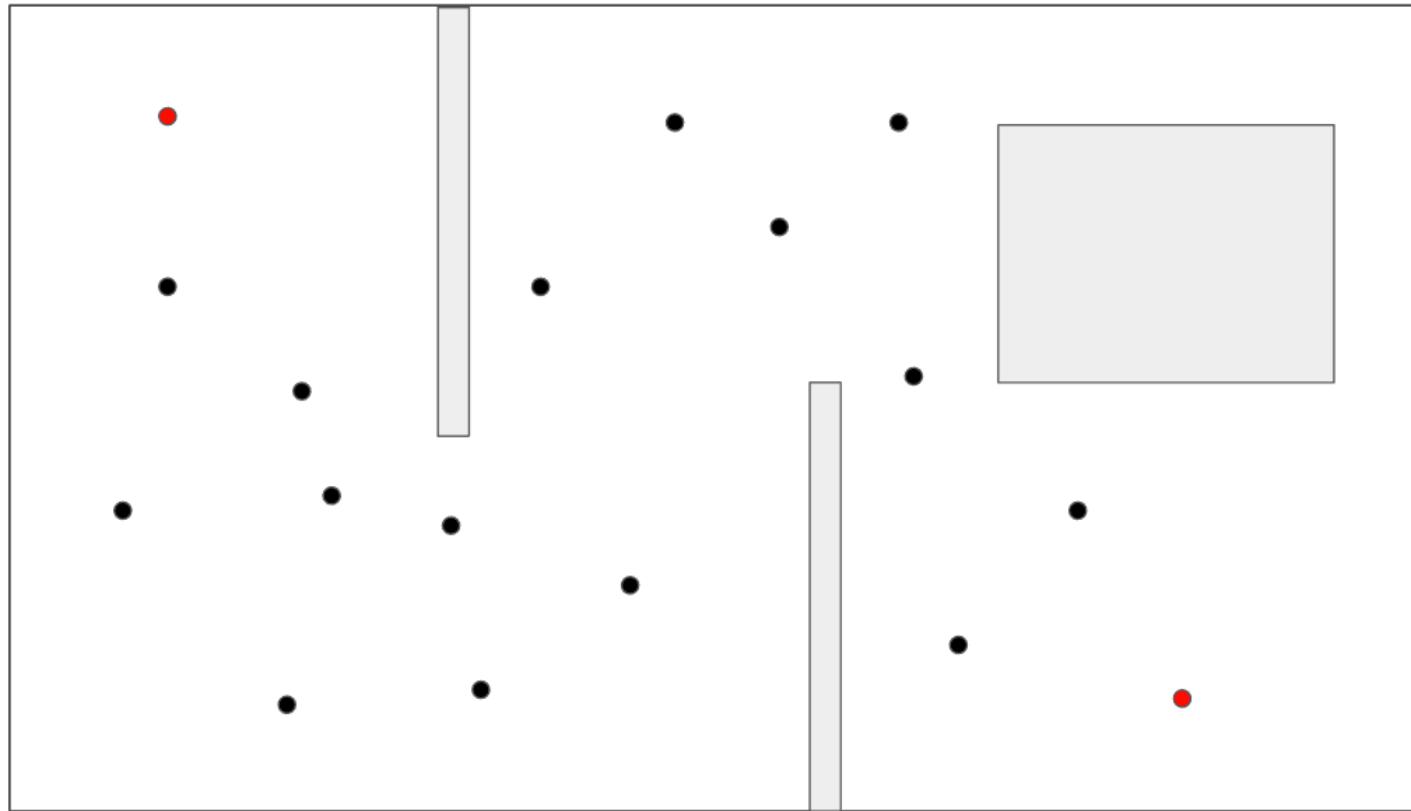
What do we do if we have no grid, or the required grid dimensions make computation infeasible, or we have many degrees of freedom?

Succinctly: Sampling-based methods are good **for large-scale, continuous environments with high-dimensional robots**

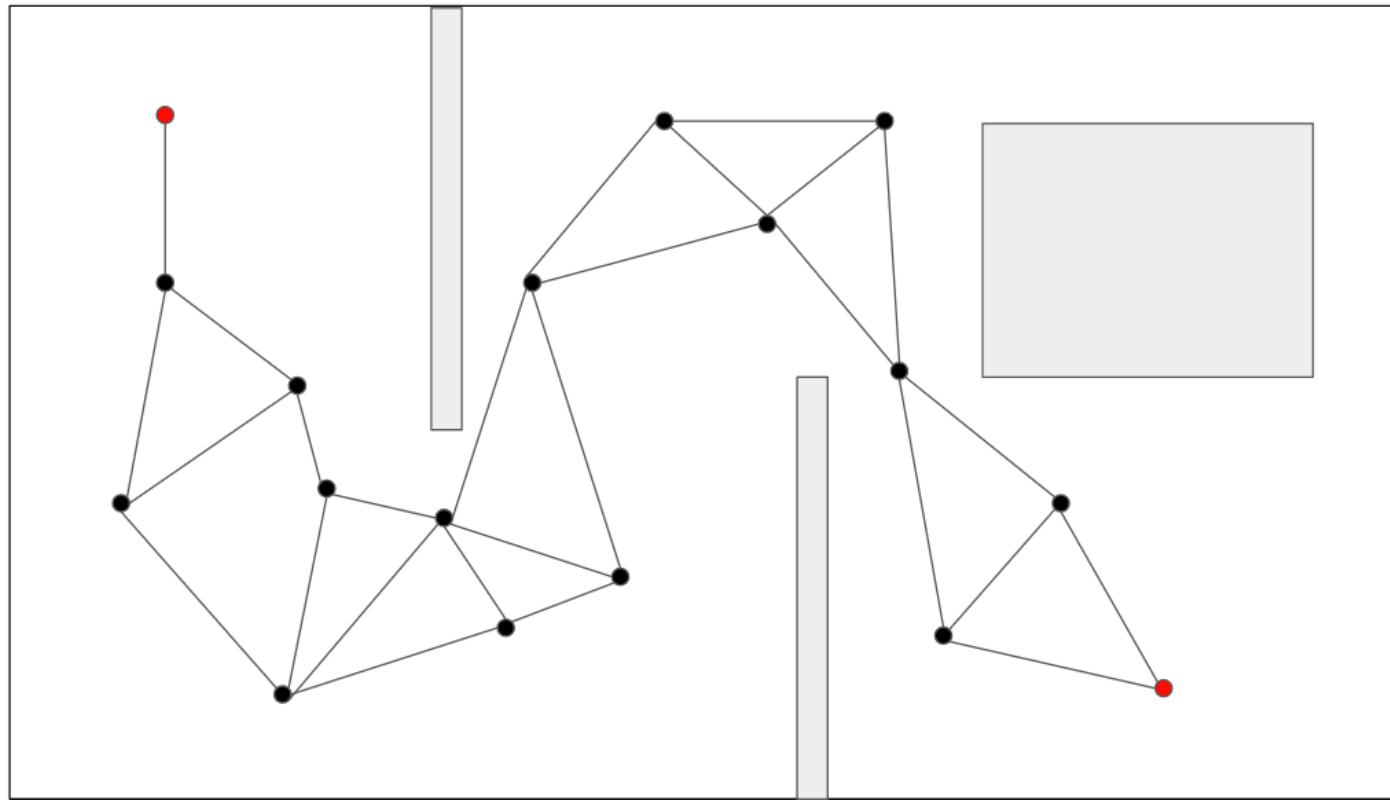
Single Robot Planning: Probabilistic Roadmap (PRM)



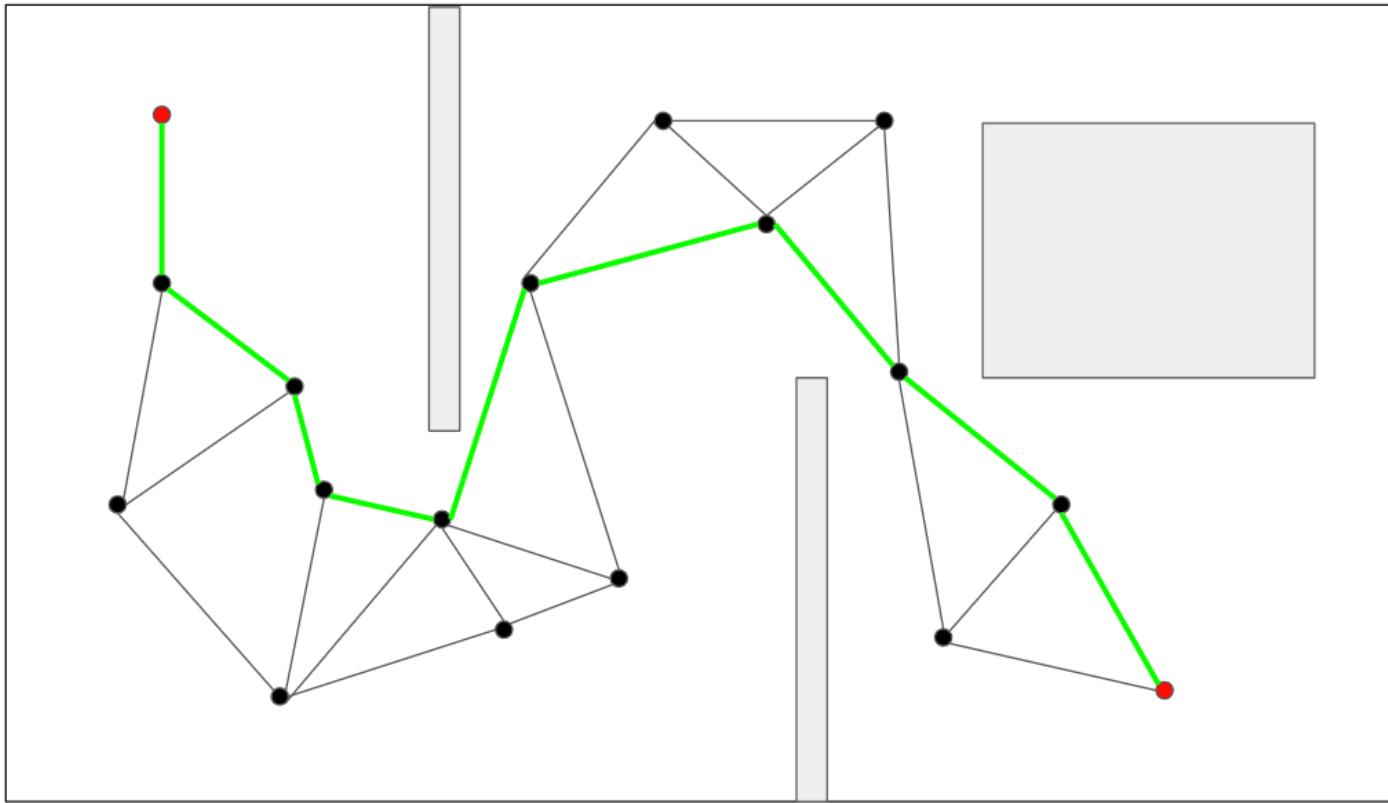
Single Robot Planning: PRM



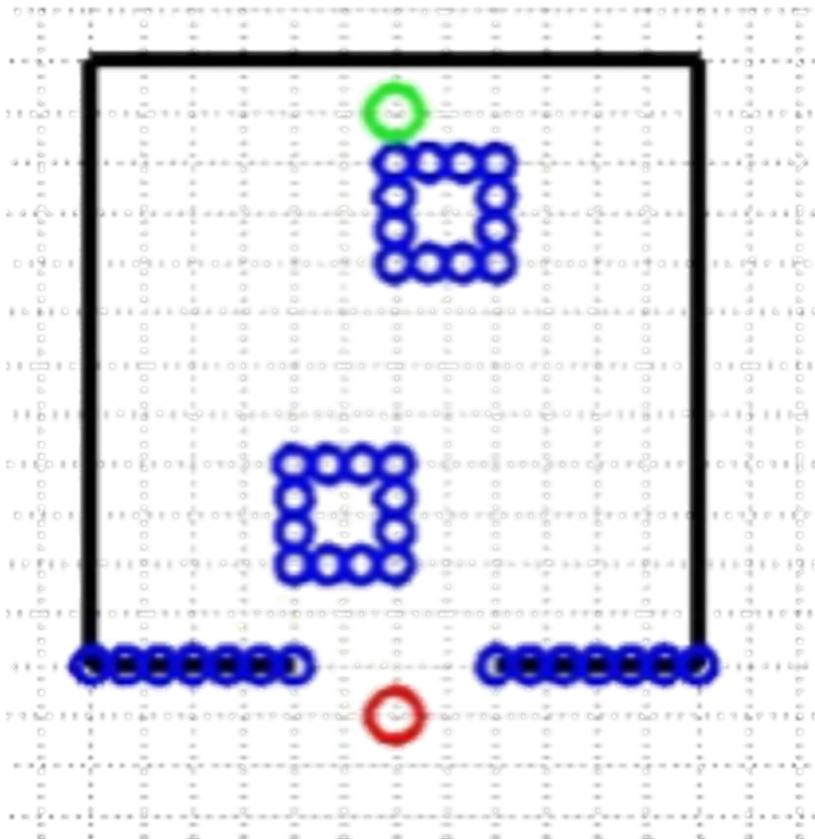
Single Robot Planning: PRM



Single Robot Planning: PRM



Single Robot Planning: RRT

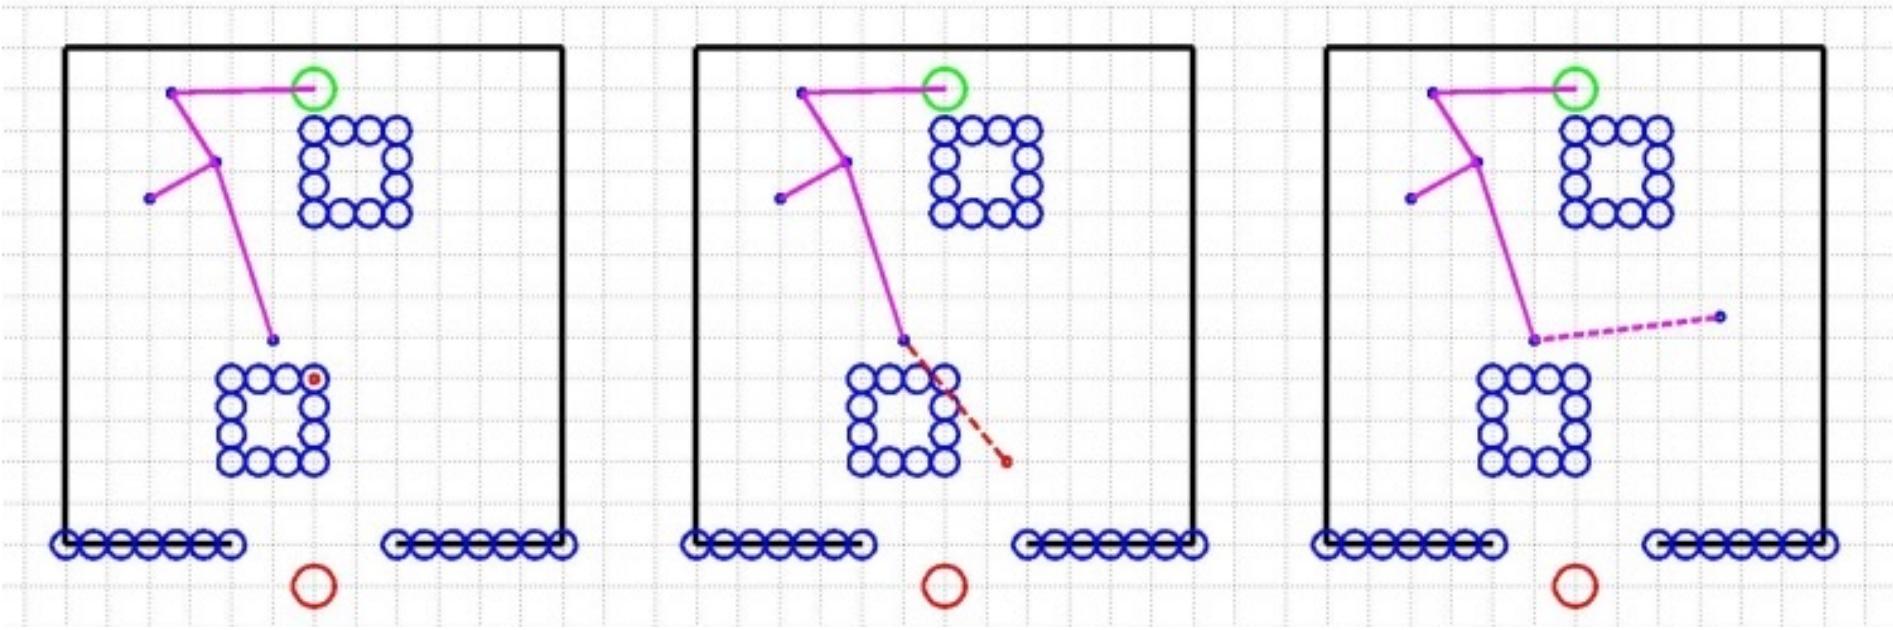


Instead of a discretized grid, positions are sampled from the continuous state space
(grid in images shown for visualization only)

Initial position $p_0 = [x_0, y_0]$

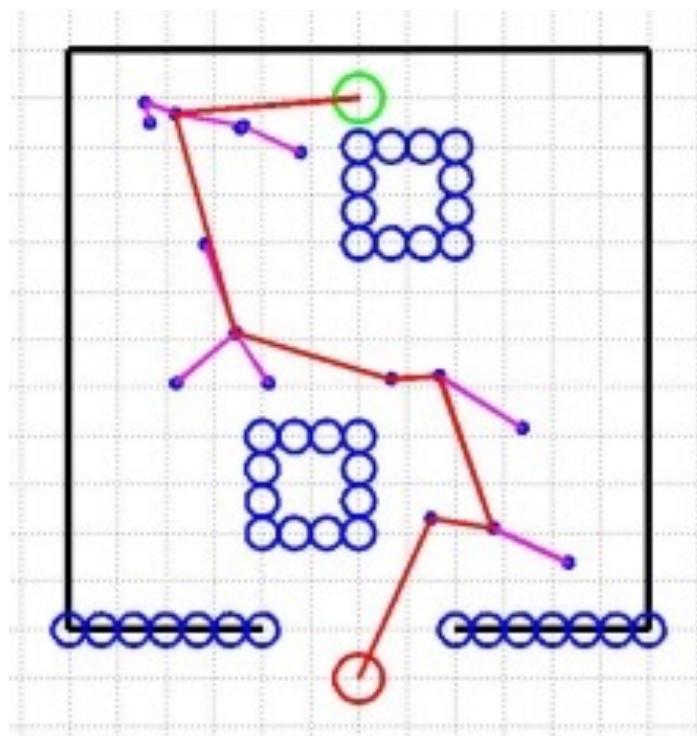
Target position $p_t = [x_t, y_t]$

Single Robot Planning: RRT

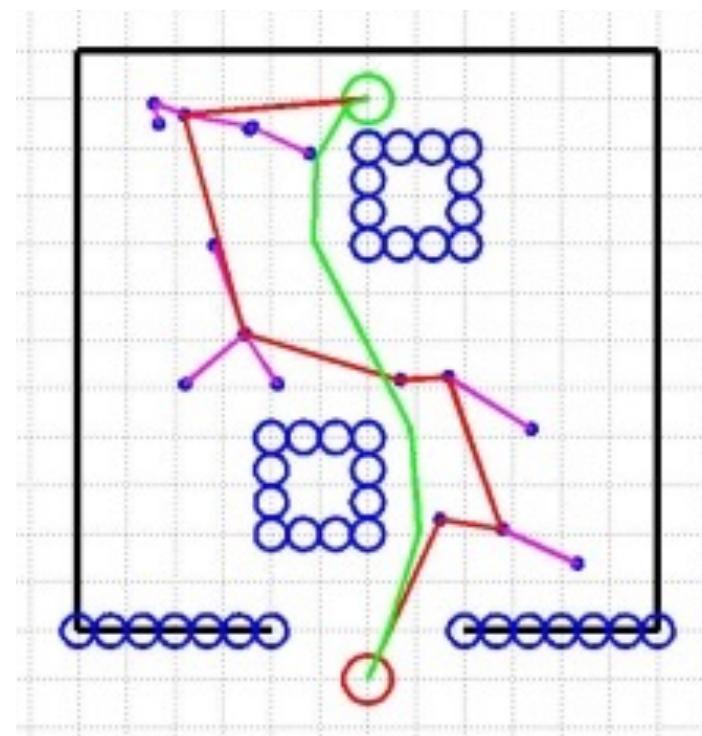


Points are added to the “tree”, then checked for 1) collisions with obstacles, 2) collisions with obstacles between nodes

Single Robot Planning: RRT



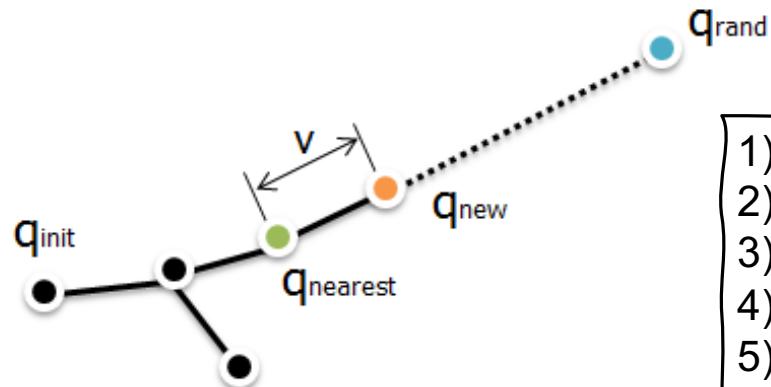
Valid path formed to target



Path can be “smoothed” by re-sampling between path points

Single Robot Planning: RRT

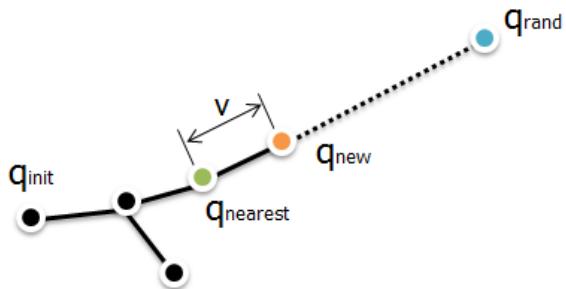
```
1  $V \leftarrow \{x_{\text{init}}\}; E \leftarrow \emptyset; i \leftarrow 0;$ 
2 while  $i < N$  do
3    $G \leftarrow (V, E);$ 
4    $x_{\text{rand}} \leftarrow \text{Sample}(i); i \leftarrow i + 1;$ 
5    $(V, E) \leftarrow \text{Extend}(G, x_{\text{rand}});$ 
```



- 1) Find the nearest node in tree
- 2) Generate a new node in direction of sampled node
- 3) Collision check
- 4) Add new node to tree
- 5) Return status (reached/advanced/trapped)

Single Robot Planning: RRT

```
1  $V \leftarrow \{x_{\text{init}}\}$ ;  $E \leftarrow \emptyset$ ;  $i \leftarrow 0$ ;  
2 while  $i < N$  do  
3    $G \leftarrow (V, E)$ ;  
4    $x_{\text{rand}} \leftarrow \text{Sample}(i)$ ;  $i \leftarrow i + 1$ ;  
5    $(V, E) \leftarrow \text{Extend}(G, x_{\text{rand}})$ ;
```



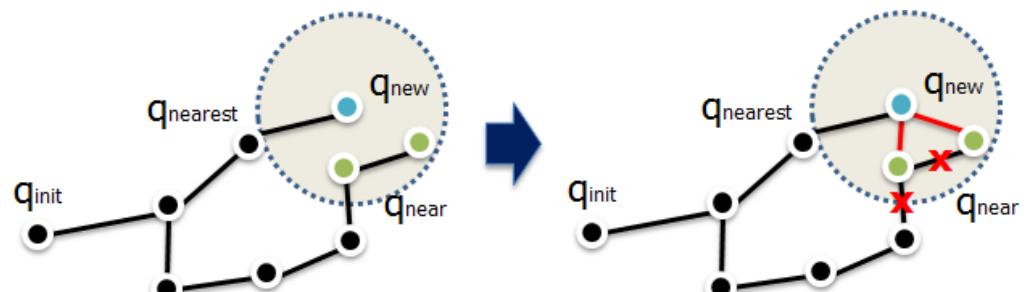
```
1  $V' \leftarrow V$ ;  $E' \leftarrow E$ ;  
2  $x_{\text{nearest}} \leftarrow \text{Nearest}(G, x)$ ;  
3  $x_{\text{new}} \leftarrow \text{Steer}(x_{\text{nearest}}, x)$ ;  
4 if  $\text{ObstacleFree}(x_{\text{nearest}}, x_{\text{new}})$  then  
5    $V' \leftarrow V' \cup \{x_{\text{new}}\}$ ;  
6    $E' \leftarrow E' \cup \{(x_{\text{nearest}}, x_{\text{new}})\}$ ;  
7 return  $G' = (V', E')$ 
```

Single Robot Planning: RRT*

```

1  $V' \leftarrow V; E' \leftarrow E;$ 
2  $x_{\text{nearest}} \leftarrow \text{Nearest}(G, x);$ 
3  $x_{\text{new}} \leftarrow \text{Steer}(x_{\text{nearest}}, x);$ 
4 if ObstacleFree( $x_{\text{nearest}}, x_{\text{new}}$ ) then
5    $V' \leftarrow V' \cup \{x_{\text{new}}\};$ 
6    $x_{\text{min}} \leftarrow x_{\text{nearest}};$ 
7    $X_{\text{near}} \leftarrow \text{Near}(G, x_{\text{new}}, |V|);$ 
8   for all  $x_{\text{near}} \in X_{\text{near}}$  do
9     if ObstacleFree( $x_{\text{near}}, x_{\text{new}}$ ) then
10       $c' \leftarrow \text{Cost}(x_{\text{near}}) + c(\text{Line}(x_{\text{near}}, x_{\text{new}}));$ 
11      if  $c' < \text{Cost}(x_{\text{new}})$  then
12         $x_{\text{min}} \leftarrow x_{\text{near}};$ 
13
14    $E' \leftarrow E' \cup \{(x_{\text{min}}, x_{\text{new}})\};$ 
15   for all  $x_{\text{near}} \in X_{\text{near}} \setminus \{x_{\text{min}}\}$  do
16     if ObstacleFree( $x_{\text{new}}, x_{\text{near}}$ ) and
17        $\text{Cost}(x_{\text{near}}) > \text{Cost}(x_{\text{new}}) + c(\text{Line}(x_{\text{new}}, x_{\text{near}}))$  then
18        $x_{\text{parent}} \leftarrow \text{Parent}(x_{\text{near}});$ 
19        $E' \leftarrow E' \setminus \{(x_{\text{parent}}, x_{\text{near}})\};$ 
20        $E' \leftarrow E' \cup \{(x_{\text{new}}, x_{\text{near}})\};$ 
21
22 return  $G' = (V', E')$ 

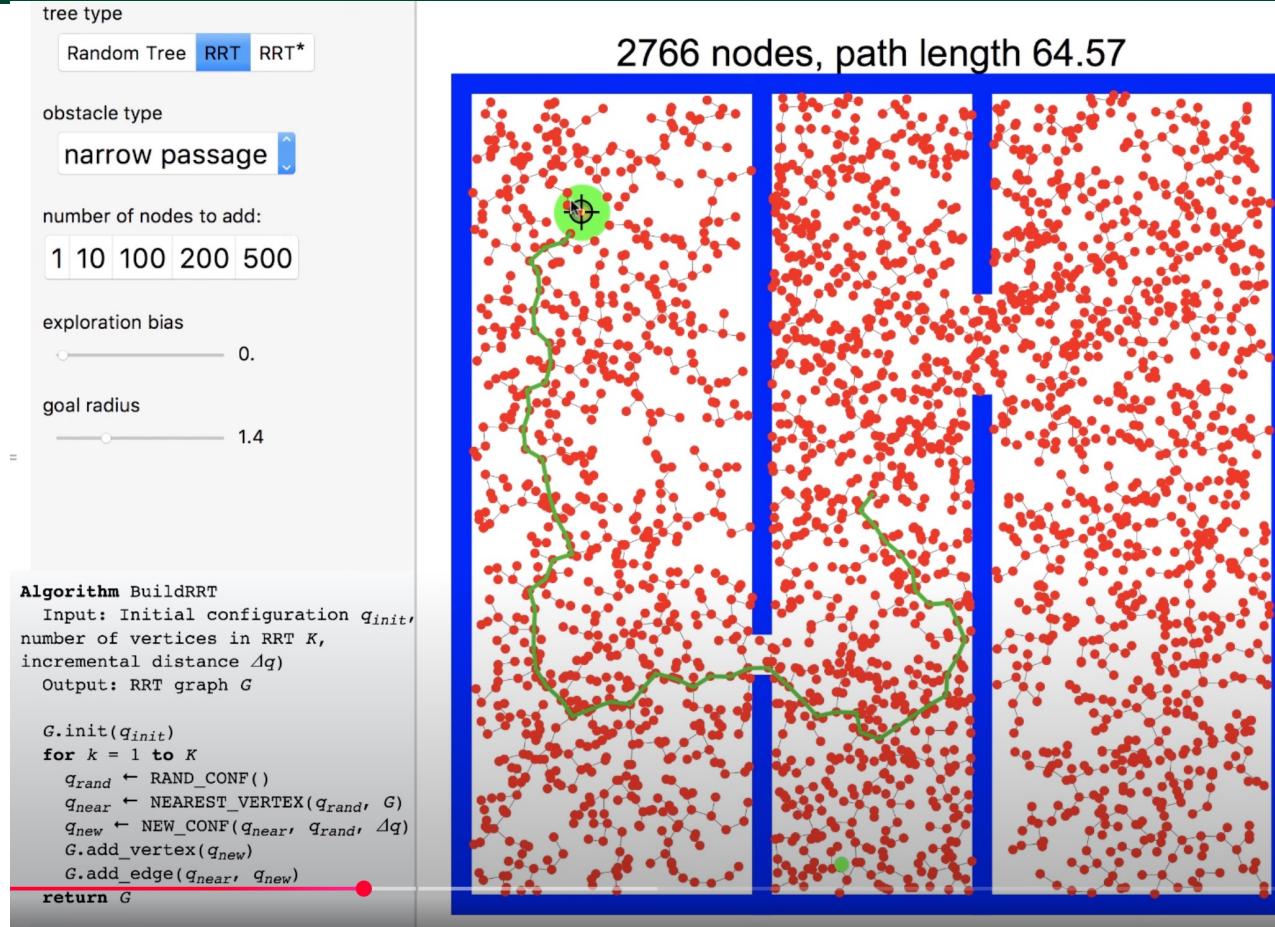
```



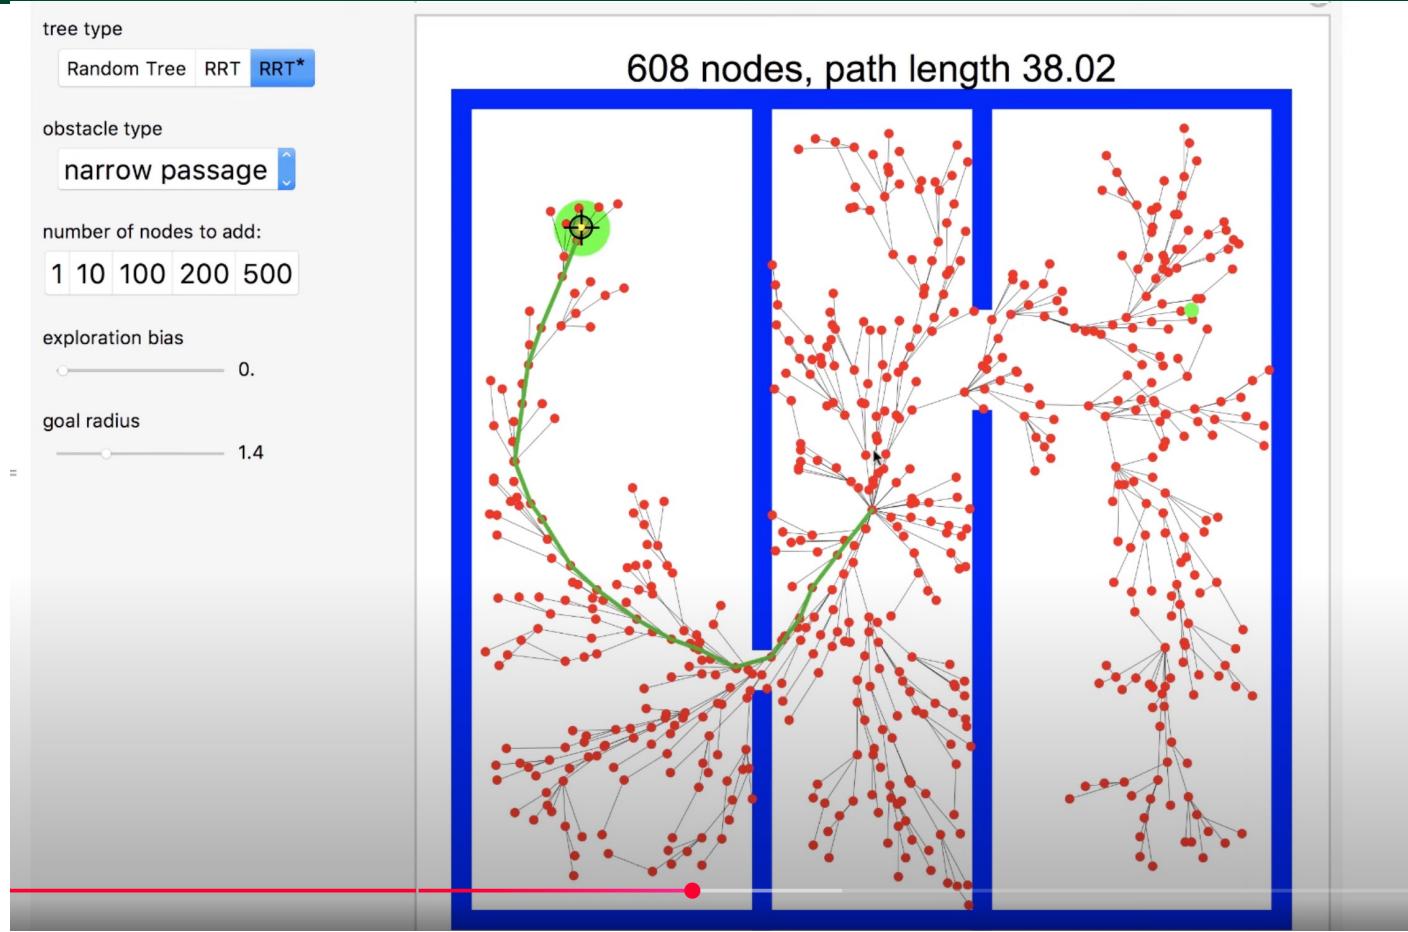
Continuously “**rewire**” the graph based on closest vertices

- 1) Find the nearest node in tree
- 2) Generate a new node in direction of sampled node
- 3) Collision check
- 4) Check whether new node is closer to any existing nodes**
- 5) Rewire tree to result in lowest total cost path**

Single Robot Planning: RRT Visualized



Single Robot Planning: RRT* Visualized



Multi-Agent RRT (MA-RRT*)

(first extend RRT* with **G-RRT***, which discretizes state space into a motion graph; when expanding from start, still samples points randomly, then uses graph-based search to find optimal path)

Then consider: **State space:** $x = (w_1, \dots, w_n)_i$ where w_1 is waypoint of agent 1;
waypoint includes time

Distance between two joint spaces: $\sum_{i=1}^n cost_i^{LB}(x_i, y_i)$

Algorithm 4 GREEDY(G_M, s, d)

```
1:  $x \leftarrow s$ ;  $c \leftarrow 0$ ;  $path \leftarrow (\emptyset, \dots, \emptyset)$ 
2: while  $x \neq d$  and  $c \leq c_{max}$  do
3:    $(path_1, \dots, path_n) \leftarrow path$ 
4:   for all  $x_i \in x$  do
5:      $N \leftarrow children(G_M, x_i)$ 
6:      $x' \leftarrow \arg \min_{x \in children(G_M, x_i)} h(x_i)$ 
7:      $c \leftarrow c + cost(x_i, x'_i)$ ;  $path_i \leftarrow path_i \cup (x_i, x'_i)$ ;
8:      $x_i \leftarrow x'_i$ 
9:   end for
10:  if not COLLISIONFREE( $path_1, \dots, path_n$ ) then
11:    return  $path$ 
12:  else
13:     $path \leftarrow (path_1, \dots, path_n)$ 
14:  end if
15: end while
16: return ( $x, path$ )
```

(obviously, this hurts our spatial intuition)

Discrete RRT (dRRT / dRRT*)

First compute *individual optimal roadmaps* using G-RRT*
THEN sample from joint configuration space of the roadmaps

Why? The joint configuration space is highly dimensional;
faster to solve individual roadmaps then only sample and
replan from those “good” solutions



Forming the solution set

For each robot:

- If H_i and U_i are empty, add to S
- If H_i is not empty, reject from S
- Otherwise, consider the cost of adding to S

$$cost(i) = |U_i|$$

Number of
robots that, if i
is added,
would be
rejected from S

- If $cost(i) \leq cost(j) \forall j \in U_i$, add to S

Example:

$$H_i = \emptyset$$

$$L_i = \emptyset$$

$$U_i = \{j\}$$

$$H_j = \emptyset$$

$$L_j = \emptyset$$

$$U_j = \{i, k\}$$

* $|U_1| = 1, |U_2| = 2$, so robot 1

will be added to solution set

Quick Summary

For scenarios with an omniscient planner: A* is ubiquitous (e.g., video games)

For highly dynamic environments: D*-lite is state of the art

For large, high-dimensional state spaces: dRRT* is state of the art

Multi-agent coordination: D*-lite does not explicitly account for other robots!

Common approach: “prioritized planning”, where high-priority robots maintain their path, low-priority robots must replan.

1. Lecture (~20 minutes)
2. Subsystem check-in (~40 minutes)

Questions?

Lecture 1: MRS Survey

Lecture 2: Motion Planning

Lecture 3: Task Assignment (+Motion Planning if needed)

Lecture 3: Communication and Control Paradigms

693H Robotic Subsystems

- Second Design Review Presentations are **3/10, 3/12**
- You will be expected to have **functional prototypes** complete to show by then
- You will be expected to have **PCB designs and final BoMs** complete by then

